

Import frames and train, test, evaluate network on facial images

✧ Importing Dependencies

```
1 from google.colab import drive
2 drive.mount('/content/drive', force_remount=True)
```

Mounted at /content/drive

```
1 PROJECT_FOLDER = '/content/drive/MyDrive/CNN_Emotion_Classification/'
```

```
1 # %cd /content/drive/MyDrive/CNN_Emotion_Classification/
```

```
1 # copy content in main folder
2 ! cp -a {PROJECT_FOLDER}. ./
```

```
1 !pip install -r Requirements/pireqs_opencv_contrib_env.txt
```

Show hidden output

Double-click (or enter) to edit

```
1 !pip install mlflow databricks-sdk
```

Requirement already satisfied: mlflow in /usr/local/lib/python3.10/dist-packages (2.13.2)
Collecting databricks-sdk
 Downloading databricks_sdk-0.28.0-py3-none-any.whl (497 kB)
 497.1/497.1 kB 5.1 MB/s eta 0:00:00

Requirement already satisfied: Flask<4 in /usr/local/lib/python3.10/dist-packages (from mlflow) (2.2.5)
Requirement already satisfied: alembic!=1.10.0,<2 in /usr/local/lib/python3.10/dist-packages (from mlflow) (1.13.1)
Requirement already satisfied: cachetools<6,>=5.0.0 in /usr/local/lib/python3.10/dist-packages (from mlflow) (5.3.3)
Requirement already satisfied: click<9,>=7.0 in /usr/local/lib/python3.10/dist-packages (from mlflow) (8.1.7)
Requirement already satisfied: cloudpickle<4 in /usr/local/lib/python3.10/dist-packages (from mlflow) (2.2.1)
Requirement already satisfied: docker<8,>=4.0.0 in /usr/local/lib/python3.10/dist-packages (from mlflow) (7.1.0)
Requirement already satisfied: entrypoints<1 in /usr/local/lib/python3.10/dist-packages (from mlflow) (0.4)
Requirement already satisfied: gitpython<4,>=3.1.9 in /usr/local/lib/python3.10/dist-packages (from mlflow) (3.1.43)
Requirement already satisfied: graphene<4 in /usr/local/lib/python3.10/dist-packages (from mlflow) (3.3)
Requirement already satisfied: importlib-metadata!=4.7.0,<8,>=3.7.0 in /usr/local/lib/python3.10/dist-packages (from mlflow) (7.1)
Requirement already satisfied: markdown<4,>=3.3 in /usr/local/lib/python3.10/dist-packages (from mlflow) (3.6)
Requirement already satisfied: matplotlib<4 in /usr/local/lib/python3.10/dist-packages (from mlflow) (3.7.1)
Requirement already satisfied: numpy<2 in /usr/local/lib/python3.10/dist-packages (from mlflow) (1.25.2)
Requirement already satisfied: opentelemetry-api<3,>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from mlflow) (1.25.0)
Requirement already satisfied: opentelemetry-sdk<3,>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from mlflow) (1.25.0)
Requirement already satisfied: packaging<25 in /usr/local/lib/python3.10/dist-packages (from mlflow) (24.1)
Requirement already satisfied: pandas<3 in /usr/local/lib/python3.10/dist-packages (from mlflow) (2.0.3)
Requirement already satisfied: protobuf<5,>=3.12.0 in /usr/local/lib/python3.10/dist-packages (from mlflow) (3.20.3)
Requirement already satisfied: pyarrow<16,>=4.0.0 in /usr/local/lib/python3.10/dist-packages (from mlflow) (14.0.2)
Requirement already satisfied: pytz<2025 in /usr/local/lib/python3.10/dist-packages (from mlflow) (2023.4)
Requirement already satisfied: pyyaml<7,>=5.1 in /usr/local/lib/python3.10/dist-packages (from mlflow) (6.0.1)
Requirement already satisfied: querystring-parser<2 in /usr/local/lib/python3.10/dist-packages (from mlflow) (1.2.4)
Requirement already satisfied: requests<3,>=2.17.3 in /usr/local/lib/python3.10/dist-packages (from mlflow) (2.31.0)
Requirement already satisfied: scikit-learn<2 in /usr/local/lib/python3.10/dist-packages (from mlflow) (1.2.2)
Requirement already satisfied: scipy<2 in /usr/local/lib/python3.10/dist-packages (from mlflow) (1.11.4)
Requirement already satisfied: sqlalchemy<3,>=1.4.0 in /usr/local/lib/python3.10/dist-packages (from mlflow) (2.0.30)
Requirement already satisfied: sqlparse<1,>=0.4.0 in /usr/local/lib/python3.10/dist-packages (from mlflow) (0.5.0)
Requirement already satisfied: Jinja2<4,>=2.11 in /usr/local/lib/python3.10/dist-packages (from mlflow) (3.1.4)
Requirement already satisfied: unicorn<23 in /usr/local/lib/python3.10/dist-packages (from mlflow) (22.0.0)
Requirement already satisfied: google-auth~=2.0 in /usr/local/lib/python3.10/dist-packages (from databricks-sdk) (2.27.0)
Requirement already satisfied: Mako in /usr/local/lib/python3.10/dist-packages (from alembic!=1.10.0,<2->mlflow) (1.3.5)
Requirement already satisfied: typing-extensions>=4 in /usr/local/lib/python3.10/dist-packages (from alembic!=1.10.0,<2->mlflow) (4.11.0)
Requirement already satisfied: urllib3>=1.26.0 in /usr/local/lib/python3.10/dist-packages (from docker<8,>=4.0.0->mlflow) (2.0.7)
Requirement already satisfied: Werkzeug>=2.2.2 in /usr/local/lib/python3.10/dist-packages (from Flask<4->mlflow) (3.0.3)
Requirement already satisfied: itsdangerous>=2.0 in /usr/local/lib/python3.10/dist-packages (from Flask<4->mlflow) (2.2.0)
Requirement already satisfied: gitdb<5,>=4.0.1 in /usr/local/lib/python3.10/dist-packages (from gitpython<4,>=3.1.9->mlflow) (4.0.10)
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from google-auth~=2.0->databricks-sdk) (0.6.1)
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.10/dist-packages (from google-auth~=2.0->databricks-sdk) (4.9)
Requirement already satisfied: graphql-core<3.3,>=3.1 in /usr/local/lib/python3.10/dist-packages (from graphene<4->mlflow) (3.2.3)
Requirement already satisfied: graphql-relay<3.3,>=3.1 in /usr/local/lib/python3.10/dist-packages (from graphene<4->mlflow) (3.2.6)
Requirement already satisfied: aniso8601<10,>=8 in /usr/local/lib/python3.10/dist-packages (from graphene<4->mlflow) (9.0.1)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.10/dist-packages (from importlib-metadata!=4.7.0,<8,>=3.7.0->mlflow) (3.17.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2<4,>=2.11->mlflow) (2.1.5)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib<4->mlflow) (1.2.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib<4->mlflow) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib<4->mlflow) (4.53.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib<4->mlflow) (1.4.5)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib<4->mlflow) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib<4->mlflow) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib<4->mlflow) (2.8.2)

Requirement already satisfied: deprecated>=1.2.6 in /usr/local/lib/python3.10/dist-packages (from opentelemetry-api<3,>=1.0.0->mlflow)
Requirement already satisfied: opentelemetry-semantic-conventions==0.46b0 in /usr/local/lib/python3.10/dist-packages (from opentelemetry-api<3,>=1.0.0->mlflow)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas<3->mlflow) (2024.1)

```
1 from display import pltDisplay
2 from pathlib import Path
3 from matplotlib.colors import ListedColormap
4 from sklearn.metrics import classification_report, roc_curve, auc, roc_auc_score
5 from sklearn.preprocessing import LabelBinarizer
6 from sklearn.model_selection import train_test_split
7 # from sorting import human_sort
8 from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, Dense, Flatten
9 from tensorflow.keras.layers import BatchNormalization, Dropout, Activation, ReLU, Softmax, GlobalAveragePooling2D
10 from tensorflow.keras.models import Model, Sequential
11 from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau
12 from tensorflow.keras.regularizers import l2
13 # import IPython.display as ipd
14
15
16 import constants as const
17 import csv
18 import cv2
19 import json
20 import logging
21 import gc
22 import matplotlib.pyplot as plt
23 import numpy as np
24 import pandas as pd
25 import random
26 import seaborn as sns
27 import tensorflow as tf
28 import subprocess
29 import tensorflow.keras.backend as K
30 import time
31 import utils
32
33 # from keras_vggface.vggface import VGGFace
34 # from tensorflow.keras.applications import EfficientNetB0
35 from keras.applications.vgg16 import VGG16
36 from keras.applications.vgg16 import preprocess_input as preprocess_imagenet
37 # from keras_vggface.utils import preprocess_input as preprocess_vggface
38
39 import mlflow
40 import gc
```

```
1 # Set mlflow parameters
2 mlflow_params = {
3     'local': False,
4     'host': '127.0.0.1',
5     'port': '8080'
6 }
7 if (mlflow_params['local']):
8     # start server
9     mlflow_server_process = subprocess.Popen(
10         f"mlflow server --host {mlflow_params['host']} --port {mlflow_params['port']}",
11         shell=False, stdout=subprocess.PIPE, stderr=subprocess.PIPE
12     )
13     mlflow.set_tracking_uri(uri="http://127.0.0.1:8080")
14     mlflow.set_experiment("Frames")
15 else:
16     # https://www.mlflow.org/docs/latest/getting-started/tracking-server-overview/index.html method 2
17     # Host: https://community.cloud.databricks.com/
18     mlflow.login()
19     mlflow.set_tracking_uri(uri="databricks")
20     mlflow.set_experiment("/Frames")
21
22 mlflow.tensorflow.autolog(
23     log_models=False,
24     log_datasets=True,
25     log_input_examples=True,
26     log_model_signatures=True,
27     checkpoint=False
28 )
```

2024/06/13 09:41:23 INFO mlflow.utils.credentials: No valid Databricks credentials found, please enter your credentials...
Databricks Host (should begin with <https://>): <https://community.cloud.databricks.com/>
Username: micasadmail@gmail.com
Password:
2024/06/13 09:42:50 INFO mlflow.utils.credentials: Successfully connected to MLflow hosted tracking server! Host: <https://community.cloud.databricks.com/>

```
1 print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
```

↗ Num GPUs Available: 1

```
1 gpus = tf.config.list_physical_devices('GPU')
2 if gpus:
3     try:
4         # Currently, memory growth needs to be the same across GPUs
5         for gpu in gpus:
6             print("Name:", gpu.name, " Type:", gpu.device_type)
7             print(tf.config.experimental.get_device_details(gpu))
8             print(tf.config.experimental.get_memory_info('GPU:0'))
9             # tf.config.experimental.set_memory_growth(gpu, True)
10        logical_gpus = tf.config.list_logical_devices('GPU')
11        print(len(gpus), "Physical GPUs,", len(logical_gpus), "Logical GPUs")
12    except RuntimeError as e:
13        # Memory growth must be set before GPUs have been initialized
14        print(e)
```

↗ Name: /physical_device:GPU:0 Type: GPU
{'compute_capability': (7, 5), 'device_name': 'Tesla T4'}
{'current': 0, 'peak': 0}
1 Physical GPUs, 1 Logical GPUs

```
1 # unzip archive in folder
2 !7z x Generated/Frames_300.zip -oGenerated/
```

↗

7-Zip [64] 16.02 : Copyright (c) 1999-2016 Igor Pavlov : 2016-05-21
p7zip Version 16.02 (locale=en_US.UTF-8,Utf16=on,HugeFiles=on,64 bits,2 CPUs Intel(R) Xeon(R) CPU @ 2.20GHz (406F0),ASM,AES-NI)

Scanning the drive for archives:
1 file, 4154172627 bytes (3962 MiB)

Extracting archive: Generated/Frames_300.zip
--
Path = Generated/Frames_300.zip
Type = zip
Physical Size = 4154172627
64-bit = +

Everything is Ok

Folders: 202
Files: 158556
Size: 4143561934
Compressed: 4154172627

```
1 utils.modules_info()
```

↗

OpenCV:
Version: 4.8.0
Tensorflow:
Version: 2.15.0

```
1 log_file = Path(const.logs_path, 'FACER.log')
2 logging.basicConfig(
3     format='%(asctime)s %(message)s',
4     filemode='a',
5     filename=log_file,
6     encoding='utf-8',
7     level=logging.INFO,
8     force=True
9 )
```

↗ Importing Dataset

```
1 data_df = pd.read_csv(Path(const.csv_path, 'dataset.csv'))
```

↗ Frame

↗ Preparing Data

✓ Dataset Creation for ML

```
1 IMG_WIDTH = IMG_HEIGHT = 224
2 IMG_CHANNELS = 3
3 SEED = 42
4 BATCH_SIZE = 16
5 VALIDATION_SPLIT = 0.2
6 EMOTIONS_LABELS = const.EMOTIONS_LABELS # RAUDESS emotion labels
```

```
1 TOTAL_ELEMENTS = const.DATASET_TOTAL_ELEMENTS
```

```
1 label_names = const.EMOTIONS_LABELS_SORTED.copy()
2 # label_names = ['happy', 'sad']
3 label_names_gender = []
4 for em in label_names:
5     label_names_gender.append(em + '_female')
6     label_names_gender.append(em + '_male')
```

✓ Dataset Creation - NEW

```
1 actors_labels = [f'{i:02d}' for i in range(1, 25)]
2 dist_idxes = {
3     '1': [slice(0, 16), slice(16, 20), slice(20, 24)],
4     '2': [slice(8, 24), slice(4, 8), slice(0, 4)],
5     '3': [slice(4, 20), slice(20, 24), slice(0, 4)]
6 }
```

```
1 # Split actors in train, validation, test
2 dist_n = 1
3 train_idxes, val_idxes, test_idxes = [actors_labels[i] for i in dist_idxes[str(dist_n)]]
4
5 print(train_idxes, val_idxes, test_idxes)
```

```
🔗 ['01', '02', '03', '04', '05', '06', '07', '08', '09', '10', '11', '12', '13', '14', '15', '16'] ['17', '18', '19', '20'] ['21', '22', '23', '24']
```

```
1 def preprocess_image(image, preprocess_type=None):
2     if (preprocess_type == 'Imagenet'):
3         image = preprocess_imagenet(image)
4     elif (preprocess_type == 'VGGFace'):
5         image = preprocess_vggface(image, version=1)
6     else:
7         image = image / 255
8
9     return image
```

```
1 def prep_external_image(filename, preprocess_type=None):
2     image = tf.io.read_file(filename)
3     image = tf.image.decode_jpeg(image, channels=IMG_CHANNELS, fancy_upscaling=False)
4     image = tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH])
5     return image
6 # return preprocess_image(image, preprocess_type)
```

```

1 def make_dataset(path, actors_idx, talk_frame=False, acted_frame=False,
2                 undersampling=False, preprocess_vgg=True, shuffle=False,
3                 gender_classes=False, sampling=1):
4
5     def parse_image(filename):
6         image = tf.io.read_file(filename)
7         image = tf.image.decode_jpeg(image, channels=IMG_CHANNELS)
8         image = tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH])
9         if (preprocess_vgg == 'Imagenet'):
10             image = preprocess_imagenet(image)
11         elif (preprocess_vgg == 'VGGFace'):
12             image = preprocess_vggface(image, version=1)
13
14         image = image / 255
15
16         return image
17
18     filenames = []
19     talk_regex = '*-01.jpg' if talk_frame else '*.jpg'
20     acted_regex = '02' if acted_frame else ''
21     gen_regex = f'*-*-*-{acted_regex}-*-*-*-{talk_regex}'
22
23     file_dict = dict()
24     file_dict_gender = dict()
25
26     for label in sorted(label_names):
27         file_dict[label] = []
28
29     for label in sorted(label_names_gender):
30         file_dict_gender[label] = []
31
32
33     for label in label_names:
34         for actor in actors_idx:
35             for file in Path(path, label, actor).glob(f'{gen_regex}'):
36                 gender = utils.get_gender(str(file))[1]
37                 lab = label + f'_{gender}'
38                 file_dict_gender[lab].append(str(file))
39                 file_dict[label].append(str(file))
40
41     # if shuffle:
42     #     for label, item in file_dict_gender.items():
43     #         logging.info(f'Label: {label}')
44     #         logging.info(f'Array len: {len(item)}')
45     #         random.Random(SEED).shuffle(item)
46
47     arr_len = [len(arr) for arr in file_dict.values()]
48
49     if undersampling:
50         filenames = [arr[:min(arr_len)] for arr in file_dict.values()]
51     else:
52         filenames = [arr for arr in file_dict.values()]
53
54     filenames = sum(filenames, [])
55
56     if shuffle:
57         random.Random(SEED).shuffle(filenames)
58
59     labels = []
60
61     if (gender_classes):
62
63         for elem in filenames:
64             cl = utils.get_class_string(str(elem))
65             gender = utils.get_gender(str(elem))[1]
66             lab = cl + f'_{gender}'
67             labels.append(label_names_gender.index(lab))
68     else:
69         labels = [
70             label_names.index(EMOTIONS_LABELS[int(utils.get_class(elem)) - 1])
71             for elem in filenames
72         ]
73
74     if (sampling < 1):
75         filenames, _, labels, _ = train_test_split(
76             filenames, labels, train_size=sampling, random_state=SEED
77         )
78
79     filenames_ds = tf.data.Dataset.from_tensor_slices(filenames)
80     labels_ds = tf.data.Dataset.from_tensor_slices(labels)
81
82     images_ds = filenames_ds.map(

```

```

83     parse_image, num_parallel_calls=tf.data.experimental.AUTOTUNE
84 )
85 ds = tf.data.Dataset.zip((images_ds, labels_ds))
86
87 return [ds, filenames]

```

```

1 sampling_rate = 1
2 gender_classes = False
3 talk_frame = True
4 acted_frame = False
5 preprocess_vgg = False # False, Imagenet or VGGFace
6
7 if (acted_frame):
8     print('L\'emozione neutral verrà rimossa')
9     if 'neutral' in label_names:
10         label_names.remove('neutral')
11     if 'neutral_female' in label_names_gender:
12         label_names_gender.remove('neutral_female')
13     if 'neutral_male' in label_names_gender:
14         label_names_gender.remove('neutral_male')
15
16 NUM_CLASSES = len(label_names_gender) if gender_classes else len(label_names)
17
18 train_ds, train_files = make_dataset(
19     const.frames_path, train_idx, talk_frame=talk_frame, acted_frame=acted_frame,
20     preprocess_vgg=preprocess_vgg, shuffle=True, gender_classes=gender_classes,
21     sampling=sampling_rate
22 )
23
24 val_ds, val_files = make_dataset(
25     const.frames_path, val_idx, talk_frame=talk_frame, acted_frame=acted_frame,
26     preprocess_vgg=preprocess_vgg, gender_classes=gender_classes,
27     sampling=sampling_rate
28 )
29
30 test_ds, test_files = make_dataset(
31     const.frames_path, test_idx, talk_frame=talk_frame, acted_frame=acted_frame,
32     preprocess_vgg=preprocess_vgg, gender_classes=gender_classes,
33     sampling=sampling_rate
34 )
35
36 labels = label_names if not gender_classes else label_names_gender

```

```

1 # train_ds = train_ds[:31882]
2 # train_files = train_files[:31882]
3
4 # val_ds = val_ds[:8426]
5 # val_files = val_files[:8426]
6
7 # test_ds = test_ds[:8301]
8 # test_files = test_files[:8301]

```

```

1 assert len(train_ds) == len(train_files), len(train_files)
2 assert len(val_ds) == len(val_files), len(val_files)
3 assert len(test_ds) == len(test_files), len(test_files)

```

```

1 train_ds_elements = len(train_ds)
2 test_ds_elements = len(test_ds)
3 val_ds_elements = len(val_ds)

```

```

1 print(f'train_ds samples: {train_ds_elements}')
2 print(f'test_ds samples: {test_ds_elements}')
3 print(f'val_ds samples: {val_ds_elements}')

```

```

➡ train_ds samples: 67906
   test_ds samples: 17994
   val_ds samples: 17642

```

✓ Build and train the model

Add operations to reduce read latency while training the model:

`ds.batch` Combines consecutive elements of the dataset into batches. The components of the resulting element will have an additional outer dimension, which will be *batch_size*

`ds.cache` Caches the elements in this dataset.

`ds.prefetch` Allows later elements to be prepared while the current element is being processed. This often improves latency and throughput, at the cost of using additional memory to store prefetched elements.


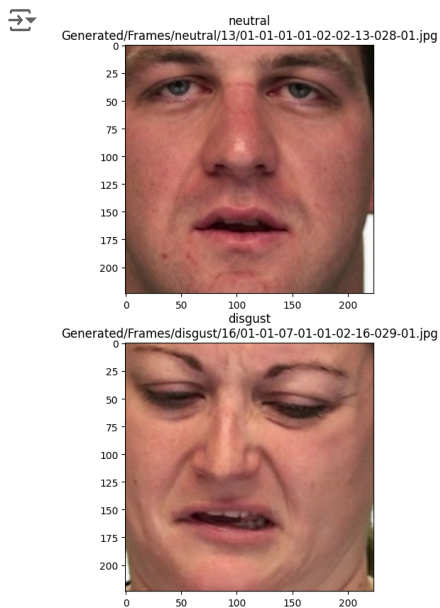
```
1 def configure_for_performance(ds, batch_size=BATCH_SIZE):
2     ds = ds.batch(batch_size)
3     # ds = ds.cache()
4     # ds = ds.shuffle(buffer_size=1000)
5     # ds = ds.repeat()
6     ds = ds.prefetch(buffer_size=tf.data.AUTOTUNE)
7     return ds
```

```
1 train_ds = configure_for_performance(train_ds)
2 val_ds = configure_for_performance(val_ds)
3 test_ds = configure_for_performance(test_ds)
```

```
1 for example_images, example_labels in train_ds.take(1):
2     print(example_images.shape)
3     print(example_labels.shape)
```

→ (16, 224, 224, 3)
(16,)

```
1 plt.figure(figsize=(16, 10))
2 rows = 2
3 cols = 2
4 n = rows * cols
5 for i in range(n):
6     plt.subplot(rows, cols, i + 1)
7     image = example_images[i]
8     plt.imshow(image) # 3 channels
9     # plt.imshow(image * 255, cmap='gray', vmin=0, vmax=255) # 1 channel
10    plt.title(f'{label_names[example_labels[i]]}\n{train_files[i]}')
```



```

1 def create_Bilotti_CNN(name='Bilotti_CNN'):
2
3     inputs = Input(shape=(IMG_HEIGHT, IMG_WIDTH, IMG_CHANNELS))
4
5     conv1 = Conv2D(32, kernel_size=(3, 3), activation='relu')(inputs)
6     conv2 = Conv2D(32, kernel_size=(3, 3), activation='relu')(conv1)
7     pool1 = MaxPooling2D(pool_size=(2, 2))(conv2)
8
9     conv3 = Conv2D(64, kernel_size=(3, 3), activation='relu')(pool1)
10    conv4 = Conv2D(64, kernel_size=(3, 3), activation='relu')(conv3)
11    pool2 = MaxPooling2D(pool_size=(2, 2))(conv4)
12
13    conv4 = Conv2D(18, kernel_size=(3, 3), activation='relu')(pool2)
14    conv5 = Conv2D(18, kernel_size=(3, 3), activation='relu')(conv4)
15    conv6 = Conv2D(18, kernel_size=(3, 3), activation='relu')(conv5)
16    pool3 = MaxPooling2D(pool_size=(2, 2))(conv6)
17
18    conv7 = Conv2D(56, kernel_size=(3, 3), activation='relu')(pool3)
19    conv8 = Conv2D(56, kernel_size=(3, 3), activation='relu')(conv7)
20    conv9 = Conv2D(56, kernel_size=(3, 3), activation='relu')(conv8)
21    pool4 = MaxPooling2D(pool_size=(2, 2))(conv9)
22
23    conv10 = Conv2D(51, kernel_size=(3, 3), activation='relu')(pool4)
24    conv11 = Conv2D(51, kernel_size=(3, 3), activation='relu')(conv10)
25    conv12 = Conv2D(51, kernel_size=(3, 3), activation='relu')(conv11)
26    pool5 = MaxPooling2D(pool_size=(2, 2))(conv12)
27
28    flatten = Flatten()(pool5)
29
30    dense1 = Dense(2048, activation='relu')(flatten)
31    drop1 = Dropout(0.25)(dense1)
32
33    dense2 = Dense(1024, activation='relu')(drop1)
34    drop2 = Dropout(0.4)(dense2)
35
36    output = Dense(NUM_CLASSES, activation='softmax')(drop2)
37
38    model = Model(inputs, output)
39
40    model._name = name
41
42    return model

```

```

1 def create_VGG16_Imagenet(name='VGG16_Imagenet'):
2
3     base_model = VGG16(
4         weights='imagenet',
5         include_top=False,
6         input_shape=(IMG_WIDTH, IMG_HEIGHT, IMG_CHANNELS)
7     )
8     base_model.trainable = False # Not trainable weights
9
10    flatten_layer = Flatten()
11    dense_layer_1 = Dense(2048, activation='relu')
12    drop_1 = Dropout(0.4)
13    dense_layer_2 = Dense(1024, activation='relu')
14    drop_2 = Dropout(0.4)
15    dense_layer_3 = Dense(512, activation='relu')
16    drop_3 = Dropout(0.4)
17    prediction_layer = Dense(NUM_CLASSES, activation='softmax')
18
19    model = Sequential([
20        base_model,
21        flatten_layer,
22        dense_layer_1,
23        drop_1,
24        dense_layer_2,
25        drop_2,
26        dense_layer_3,
27        drop_3,
28        prediction_layer
29    ])
30
31    model._name = name
32
33    return model

```



```

1 def create_EfficientNetB0_Imagenet(name='EfficientNetB0_Imagenet'):
2
3     nb_class = NUM_CLASSES
4
5     inputs = Input(shape=(IMG_WIDTH, IMG_HEIGHT, IMG_CHANNELS))
6     model = EfficientNetB0(
7         include_top=False, input_tensor=inputs, weights="imagenet"
8     )
9
10    # Freeze the pretrained weights
11    model.trainable = False
12
13    # Rebuild top
14    x = GlobalAveragePooling2D(name="avg_pool")(model.output)
15    x = BatchNormalization()(x)
16
17    top_dropout_rate = 0.2
18    x = Dropout(top_dropout_rate, name="top_dropout")(x)
19    outputs = Dense(nb_class, activation="softmax", name="pred")(x)
20
21    model._name = name
22    # Compile
23    model = Model(inputs, outputs, name="EfficientNet")
24
25    return model

```

```

1 def create_VGG16_VGGFACE(name='VGG16_VGGFACE'):
2     nb_class = NUM_CLASSES
3
4     vgg_model = VGGFace(
5         include_top=False, weights='vggface', input_shape=(IMG_WIDTH, IMG_HEIGHT, IMG_CHANNELS)
6     )
7     last_layer = vgg_model.get_layer('pool5').output
8     x = Flatten(name='flatten')(last_layer)
9
10    x = Dense(512, activation='relu', name='fc6')(x)
11    x = Dropout(0.35)(x)
12    x = Dense(256, activation='relu', name='fc7')(x)
13    x = Dropout(0.35)(x)
14    x = Dense(128, activation='relu', name='fc8')(x)
15    x = Dropout(0.35)(x)
16
17    out = Dense(nb_class, activation='softmax', name='fc9')(x)
18
19    custom_vgg_model = Model(vgg_model.input, out)
20    custom_vgg_model._name = name
21
22    return custom_vgg_model

```

```
1 NUM_CLASSES
```

↔ 8

```

1 def create_dog_vs_cat_model(name='dog_cat'):
2     vggmodel = VGG16(weights='imagenet', include_top=True)
3     for layers in (vggmodel.layers)[:19]:
4         layers.trainable = True
5
6     X = vggmodel.layers[-2].output
7     # X = Dense(256, activation='relu')(X)
8     # predictions = Dense(1, activation="sigmoid")(X)
9     predictions = Dense(NUM_CLASSES, activation="softmax")(X)
10    model_final = Model(vggmodel.input, predictions)
11
12    model_final._name = name
13
14    return model_final

```

```
1 tf.keras.backend.clear_session() # clear all precedent models and sessions
```

```


1 check_path = 'checkpoint.weights.h5'
2 checkpointer = ModelCheckpoint(
3     check_path, monitor='val_accuracy', verbose=1, save_best_only=True,
4     save_weights_only=False, mode='auto', save_freq='epoch'
5 )

```

```

1 # model = create_cnn_model()
2 # model = medium_model()
3 # model = create_VGG16_Imagenet()
4 # model = create_VGG16_VGGFACE()
5 # model = create_EfficientNetB0_Imagenet()
6 # model = create_grigorasi_model()
7 # model = create_Bilotti_CNN()
8 model = create_dog_vs_cat_model()
9 model.summary()

```

 Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels
553467096/553467096 [=====] - 6s 0us/step
Model: "dog_cat"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 4096)	102764544
fc2 (Dense)	(None, 4096)	16781312
dense (Dense)	(None, 8)	32776

=====

Total params: 134293320 (512.29 MB)

Trainable params: 134293320 (512.29 MB)

Non-trainable params: 0 (0.00 Byte)

```

1 # Define training callbacks
2
3 class TimeHistory(tf.keras.callbacks.Callback):
4     def on_train_begin(self, logs={}):
5         self.times = []
6
7     def on_epoch_begin(self, batch, logs={}):
8         self.epoch_time_start = time.time()
9
10    def on_epoch_end(self, batch, logs={}):
11        self.times.append(time.time() - self.epoch_time_start)
12
13
14 early_stopping_callback = tf.keras.callbacks.EarlyStopping(
15     monitor='val_loss',
16     verbose=1,
17     patience=3,
18     restore_best_weights=True
19 )
20
21 reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, verbose=1,
22                               patience=7, min_lr=0)

```

```
1 METRICS = ['accuracy']
```

```

1 # Model compile dog vs cat
2 from keras import optimizers
3 model.compile(
4     loss=tf.keras.losses.BinaryCrossentropy(),
5     optimizer=optimizers.SGD(learning_rate=0.00001, momentum=0),
6     metrics=["accuracy"]
7 )

```

```

1 model.compile(
2     optimizer=tf.keras.optimizers.Adam(
3         learning_rate=1e-4
4     ),
5     # optimizer=tf.keras.optimizers.SGD(), # for VGG16_VGGFACE
6     loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
7     metrics=METRICS,
8 )

```

▼ Train Model

```

1 EPOCHS = 200
2 time_callback = TimeHistory()
3 history = model.fit(
4     x=train_ds,
5     validation_data=val_ds,
6     epochs=EPOCHS,
7     callbacks=[time_callback, early_stopping_callback, checkpointer]
8 )

```

```

➡ 2024/06/13 09:49:15 INFO mlflow.utils.autologging_utils: Created MLflow autologging r
2024/06/13 09:51:38 INFO mlflow.types.utils: MLflow 2.9.0 introduces model signature
2024/06/13 09:52:04 INFO mlflow.types.utils: MLflow 2.9.0 introduces model signature
Epoch 1/200
4245/4245 [=====] - ETA: 0s - loss: 0.4982 - accuracy: 0.818
Epoch 1: val_accuracy improved from -inf to 0.51581, saving model to checkpoint.weigh
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103: UserWarnin
saving_api.save_model(
4245/4245 [=====] - 1183s 274ms/step - loss: 0.4982 - accura
Epoch 2/200
4245/4245 [=====] - ETA: 0s - loss: 0.0885 - accuracy: 0.976
Epoch 2: val_accuracy improved from 0.51581 to 0.54535, saving model to checkpoint.we
4245/4245 [=====] - 1146s 270ms/step - loss: 0.0885 - accura
Epoch 3/200
4245/4245 [=====] - ETA: 0s - loss: 0.0520 - accuracy: 0.983
Epoch 3: val_accuracy did not improve from 0.54535
4245/4245 [=====] - 1135s 267ms/step - loss: 0.0520 - accura
Epoch 4/200
4245/4245 [=====] - ETA: 0s - loss: 0.0399 - accuracy: 0.987

Epoch 4: val_accuracy did not improve from 0.54535
4245/4245 [=====] - 1134s 267ms/step - loss: 0.0399 - accura
Epoch 4: early stopping
Uploading artifacts: 100%
2/2 [00:00<00:00, 2.03it/s]

```

```

1 # # FINE TUNING
2 # # unfreeze base model
3 # vgg_model.trainable = True
4 # model.summary()
5
6 # model.compile(
7 #     optimizer=tf.keras.optimizers.Adam(
8 #         learning_rate=1e-5
9 #     ),
10 #     # optimizer=tf.keras.optimizers.SGD(),
11 #     loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
12 #     metrics=METRICS,
13 # )
14 # EPOCHS = 2
15 # time_callback = TimeHistory()
16 # history = model.fit(
17 #     x=train_ds,
18 #     validation_data=val_ds,
19 #     epochs=EPOCHS,
20 #     callbacks=[time_callback, early_stopping_callback, reduce_lr]
21 # )

```

```

1 EPOCHS = len(time_callback.times)

```

```

1 # Create model path
2 model_path = Path(const.models_path, 'Frame', model._name)
3 run_folders = list(Path(const.models_path, 'Frame', model._name).glob('Run_*'))
4
5 if not run_folders:
6     model_path = Path(model_path, 'Run_1')
7 else:
8     last_run = run_folders.pop()
9     last_run_idx = Path(last_run).name.split('_')[-1]
10    model_path = Path(model_path, f'Run_{int(last_run_idx) + 1}')
11
12 model_path.mkdir(parents=True, exist_ok=False)

```

```

1 # Save info on the indexes used for train, val and test
2 ds_info_path = Path(model_path, f'{model._name}_dataset.txt')
3 with open(ds_info_path, 'w+', newline='') as res_file:
4     res_file.write(f'Train indexes: {train_idxs}\n')
5     res_file.write(f'Train files: {train_ds_elements}\n')
6     res_file.write(f'Val indexes: {val_idxs}\n')
7     res_file.write(f'Val files: {val_ds_elements}\n')
8     res_file.write(f'Test indexes: {test_idxs}\n')
9     res_file.write(f'Test files: {test_ds_elements}\n')

```

```

1 metrics = history.history

```

```


1 mod_loss = metrics['loss']
2 mod_val_loss = metrics['val_loss']
3 mod_accuracy = metrics['accuracy']
4 mod_val_accuracy = metrics['val_accuracy']
5 # mod_f1 = metrics['fBeta_score']
6 # mod_val_f1 = metrics['val_fBeta_score']
7
8 mod_mean_loss = np.mean(mod_loss)
9 mod_mean_val_loss = np.mean(mod_val_loss)
10 mod_mean_accuracy = np.mean(mod_accuracy)
11 mod_mean_val_accuracy = np.mean(mod_val_accuracy)
12 # mod_mean_f1 = np.mean(mod_f1)
13 # mod_mean_val_f1 = np.mean(mod_val_f1)

```


```

1 # Save Loss
2 plt.title('Model loss')
3 plt.plot(history.epoch, mod_loss, mod_val_loss)
4 plt.legend(['loss', 'val_loss'])
5 plt.ylim([0, max(plt.ylim())])
6 plt.xlabel('Epoch')
7 plt.ylabel('Loss [CrossEntropy]')
8 plt.savefig(Path(model_path, 'loss.png'))

```



```
1 # Save Accuracy
2 plt.title('Model accuracy')
3 plt.plot(
4     history.epoch,
5     100 * np.array(mod_accuracy),
6     100 * np.array(mod_val_accuracy)
7 )
8 plt.legend(['accuracy', 'val_accuracy'])
9 plt.ylim([0, 100])
10 plt.xlabel('Epoch')
11 plt.ylabel('Accuracy [%]')
12 plt.savefig(Path(model_path, 'accuracy.png'))
```



Unsupported Cell Type. Double-Click to inspect/edit the content.

Unsupported Cell Type. Double-Click to inspect/edit the content.

✓ Evaluate Model

```
1 model_eval = model.evaluate(test_ds, return_dict=True)
```



1125/1125 [=====] - 80s 71ms/step - loss: 2.2596 - accuracy: 0.4440

```

1 hist_df = pd.DataFrame(metrics)
2 hist_csv_file = Path(model_path, 'history.csv')
3 with open(hist_csv_file, mode='w') as f:
4     hist_df.to_csv(f)

```

```

1 # Save model
2 model.save(Path(model_path, f'{model._name}.keras'), overwrite=False)
3 # Save history
4 # np.save(Path(model_path, f'{model._name}_history.npy'), history)
5 # Save model image
6 model_img = tf.keras.utils.plot_model(
7     model, Path(model_path, f'{model._name}.png'), show_shapes=True,
8     show_layer_names=True, show_layer_activations=True
9 )

```

```
1 model_eval
```

```
↩ {'loss': 2.259594678878784, 'accuracy': 0.4439813196659088}
```

```

1 test_loss = model_eval['loss']
2 test_accuracy = model_eval['accuracy']
3 # test_f1 = model_eval['fBeta_score']
4 mean_epoch_time = np.mean(time_callback.times)

```

```

1 # Salvataggio informazioni modello
2 model_save_path = Path(model_path, f'{model._name}_result.txt')
3 with open(model_save_path, 'w+', newline='') as res_file:
4     res_file.write(f'BATCH: {BATCH_SIZE}\n')
5     res_file.write(f'Train loss: {str(mod_loss)}\n')
6     res_file.write(f'val_loss: {str(mod_val_loss)}\n')
7     res_file.write(f'Train accuracy: {str(mod_accuracy)}\n')
8     res_file.write(f'Train val_accuracy: {str(mod_val_accuracy)}\n')
9     # res_file.write(f'Train f1_score: {str(mod_f1)}\n')
10    # res_file.write(f'Train val_f1_score: {str(mod_val_f1)}\n')
11    res_file.write(f'Test loss: {str(test_loss)}\n')
12    res_file.write(f'Test accuracy: {str(test_accuracy)}\n')
13    # res_file.write(f'Test f1_score: {str(test_f1)}\n')
14    res_file.write(f'Mean epoch time: {str(mean_epoch_time)}')

```

```

1 # Salvataggio informazioni generali modelli
2 with open(Path(const.models_path, 'Frame', 'models.csv'), 'a+') as csvfile:
3     filewriter = csv.writer(
4         csvfile, delimiter=';', quotechar='|', quoting=csv.QUOTE_MINIMAL
5     )
6
7     # filewriter.writerow(
8     #     ["Model Name", "Epochs", "% Validation", "% Test set",
9     #     "Train loss", "Train accuracy", "Val loss", "Val accuracy",
10    #     "Test loss", "Test accuracy", "Mean epoch time", "Note"]
11    # )
12    test_ds_perc = utils.trunc((test_ds_elements * 100) / TOTAL_ELEMENTS, 2)
13    val_ds_perc = utils.trunc((val_ds_elements * 100) / TOTAL_ELEMENTS, 2)
14    full_path = str(Path(model._name, model_path.name))
15    filewriter.writerow(
16        [full_path, EPOCHS, val_ds_perc, test_ds_perc,
17         mod_loss, mod_accuracy, mod_val_loss, mod_val_accuracy,
18         test_loss, test_accuracy, mean_epoch_time, '']
19    )

```

```

1 for test_images, test_labels in test_ds.take(1):
2     print(test_images.shape)
3     print(test_labels.shape)

```

```
↩ (16, 224, 224, 3)
(16,)
```

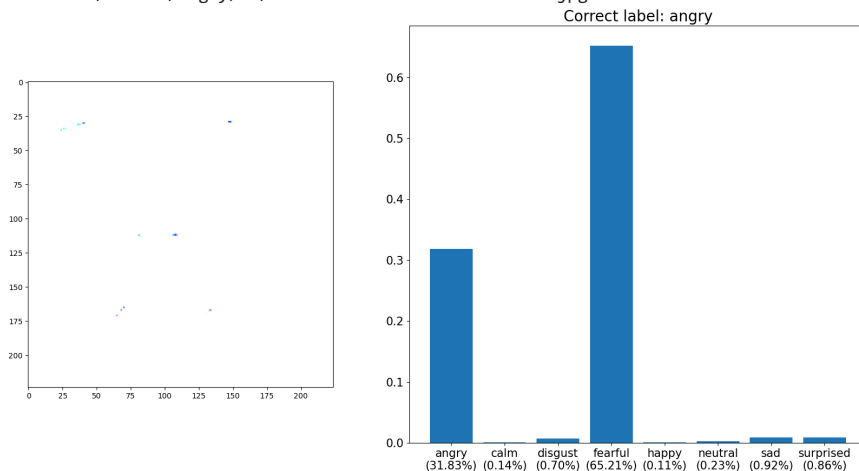
```

1 gen = np.random.default_rng(seed=None)
2 idx = gen.integers(0, len(test_images))
3 print(test_files[idx])
4
5 image = test_images[idx]
6 label = test_labels[idx]
7
8 net_input = utils.extend_tensor(image, 0)
9 prediction = model(net_input)
10 prediction = prediction[0].numpy()
11
12 valued_arr = []
13
14 for idx, name in enumerate(label_names):
15     valued_arr.append(f'{name}\n({prediction[idx]:.2%})')
16
17 fig, ax = plt.subplots(
18     nrows=1, ncols=2, width_ratios=[0.4, 0.6], figsize=(20, 10)
19 )
20
21 pltDisplay(image * 255, ax=ax[0]) # 1 channel
22 # pltDisplay(image, ax=ax[0])
23
24 ax[1].bar(valued_arr, prediction)
25 plt.xticks(fontsize=15)
26 plt.yticks(fontsize=15)
27 plt.title(f'Correct label: {label_names[label]}', fontsize=20)
28 # plt.xlabel('Predicted class')
29 # plt.ylabel('Percentage')
30 plt.show()

```



Generated/Frames/angry/21/01-01-05-01-02-01-21-062-01.jpg



Display a confusion matrix

Use a [confusion matrix](#) to check how well the model did classifying each of the commands in the test set:

Unsupported Cell Type. Double-Click to inspect/edit the content.

```
1 y_pred = model.predict(test_ds)
```



1125/1125 [=====] - 80s 70ms/step

```
1 y_pred = tf.argmax(y_pred, axis=1, output_type=tf.int32)
```

```
1 y_true = tf.concat(list(test_ds.map(lambda _, lab: lab)), axis=0)
```


```
1 label_names
```

```
↩ ['angry', 'calm', 'disgust', 'fearful', 'happy', 'neutral', 'sad', 'surprised']
```

```
1 report = classification_report(  
2     y_true, y_pred, target_names=label_names,  
3     output_dict=True, zero_division='warn'  
4 )
```

```
1 rep_to_csv = pd.DataFrame(data=report).transpose()
```

```
1 fig, ax = plt.subplots(figsize=(10, 10))  
2 ax.xaxis.tick_top()  
3 sns.heatmap(rep_to_csv.iloc[:NUM_CLASSES, :3],  
4             cbar=True,  
5             square=False,  
6             annot=True,  
7             annot_kws={'size': 15},  
8             fmt='.2g',  
9             linewidths=0.5)  
10 plt.savefig(Path(model_path, f'{model._name}_f1_score.png'))  
11 plt.show()  
12  
13  
14 with sns.axes_style('white'):  
15     fig, ax = plt.subplots(figsize=(3, 5))  
16     ax.xaxis.tick_top()  
17     sns.heatmap(rep_to_csv.iloc[:NUM_CLASSES, 3:],  
18               cbar=False,  
19               square=False,  
20               annot=True,  
21               annot_kws={'size': 15},  
22               fmt='.4g',  
23               cmap=ListedColormap([('purple'])),  
24               linewidths=0.5)  
25 plt.savefig(Path(model_path, f'{model._name}_support.png'))  
26 plt.show()
```

support	
angry	2502
calm	2540
disgust	2557
fearful	2288
happy	2224
neutral	1142

```
1 report_save_path = Path(model_path, f'{model._name}_report.csv')
2 rep_to_csv.to_csv(report_save_path)
```

Unsupported Cell Type. Double-Click to inspect/edit the content.

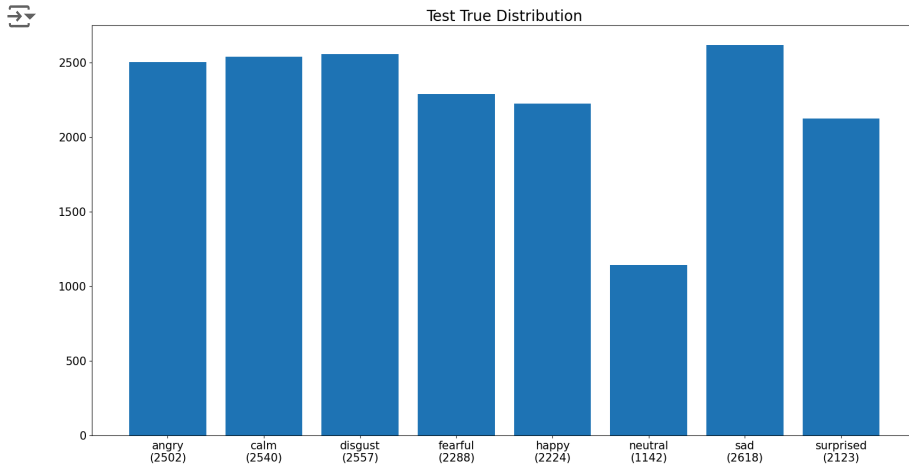
```
1 np.mean([report[c]['f1-score'] for c in list(report)[:NUM_CLASSES]])
```

0.4322732858116839

```

1 # True Test Distribution
2 # unique, counts = np.unique(y_true, return_counts=True)
3 # collections.Counter(y_true)
4 counts = [np.count_nonzero(y_true == idx) for idx in range(len(label_names))]
5 valued_arr = []
6 # for i in range(len(label_names)):
7 for idx, name in enumerate(label_names):
8     count = counts[idx]
9     valued_arr.append(f'{name}\n({count})')
10
11 fig = plt.subplots(figsize=(20, 10))
12 plt.bar(valued_arr, counts)
13 plt.xticks(fontsize=15)
14 plt.yticks(fontsize=15)
15 plt.title('Test True Distribution', fontsize=20)
16 # plt.xlabel('Predicted class')
17 # plt.ylabel('Percentage')
18 plt.savefig(Path(model_path, f'{model._name}_trueDist.png'))
19 plt.show()

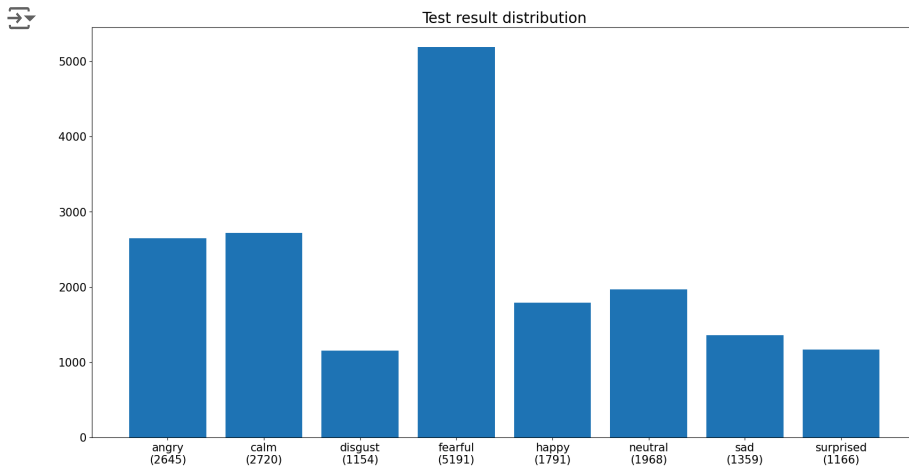
```




```

1 # Predicted Test Distribution
2 # unique, counts = np.unique(y_pred, return_counts=True)
3 counts = [np.count_nonzero(y_pred == idx) for idx in range(len(label_names))]
4 valued_arr = []
5 unique_idx = 0
6 for idx, name in enumerate(label_names):
7     count = counts[idx]
8     valued_arr.append(f'{name}\n({count})')
9
10 fig = plt.subplots(figsize=(20, 10))
11 plt.bar(valued_arr, counts)
12 plt.xticks(fontsize=15)
13 plt.yticks(fontsize=15)
14 plt.title('Test result distribution', fontsize=20)
15 # plt.xlabel('Predicted class')
16 # plt.ylabel('Percentage')
17 plt.savefig(Path(model_path, f'{model._name}_predDist.png'))
18 plt.show()

```



```
1 confusion_mtx = tf.math.confusion_matrix(y_true, y_pred)
2 fig, ax = plt.subplots(figsize=(10, 10))
3 sns.heatmap(confusion_mtx,
4             xticklabels=label_names,
5             yticklabels=label_names,
6             annot=True, fmt='.4g')
7 plt.xlabel('Prediction')
8 plt.ylabel('True Label')
9 plt.savefig(Path(model_path, f'{model._name}_heat.png'))
10 plt.show()
```



```

1 def multiclass_roc_auc_score(target, y_test, y_pred, average="macro"):
2     # function for scoring roc auc score for multi-class
3     lb = LabelBinarizer()
4     lb.fit(y_test)
5     y_test = lb.transform(y_test)
6     y_pred = lb.transform(y_pred)
7
8     if len(target) > 2:
9         for (idx, c_label) in enumerate(target):
10
11             fpr, tpr, thresholds = roc_curve(
12                 y_test[:, idx].astype(int),
13                 y_pred[:, idx]
14             )
15             c_ax.plot(
16                 fpr, tpr, label='%s (AUC:%0.2f)' % (c_label, auc(fpr, tpr))
17             )
18     else:
19         fpr, tpr, thresholds = roc_curve(
20             y_test,
21             y_pred
22         )
23         c_ax.plot(
24             fpr, tpr, label='Model (AUC:%0.2f)' % (auc(fpr, tpr)), color='#ff7f0e'
25         )
26     c_ax.plot(fpr, fpr, color='b', linestyle='--', label='Random Guessing')
27
28     return roc_auc_score(y_test, y_pred, average=average)

```

```

1 # set plot figure size
2 fig, c_ax = plt.subplots(1, 1, figsize=(12, 8))
3
4 print('ROC AUC score:', multiclass_roc_auc_score(
5     label_names,
6     tf.reshape(y_true, (y_true.shape[0], 1)),
7     tf.reshape(y_pred, (y_pred.shape[0], 1))


```

```

8 ))
9
10 c_ax.legend()
11 c_ax.set_xlabel('False Positive Rate')
12 c_ax.set_ylabel('True Positive Rate')
13 plt.savefig(Path(model_path, f'{model._name}_ROC.png'))
14 plt.show()

```

ROC AUC score: 0.6879024625140258



Run inference on an image file

Finally, verify the model's prediction output using an image

1 labels

```

['angry', 'calm', 'disgust', 'fearful', 'happy', 'sad', 'surprised']



```

```

1 # Take the first item of each class in test_files
2 # path = Path(next((subs for subs in test_files if 'angry' in subs), None))
3 # path = Path(next((subs for subs in test_files if 'calm' in subs), None))
4 # path = Path(next((subs for subs in test_files if 'disgust' in subs), None))
5 path = Path(next((subs for subs in test_files if 'fearful' in subs), None))
6 # path = Path(next((subs for subs in test_files if 'happy' in subs), None))
7 # path = Path(next((subs for subs in test_files if 'neutral' in subs), None))
8 # path = Path(next((subs for subs in test_files if 'sad' in subs), None))
9 # path = Path(next((subs for subs in test_files if 'surprised' in subs), None))
10 # path = Path('Generated/Frames/disgust/01-01-07-01-01-01-16-049-01.jpg')
11 print(path)
12 # print(tf.io.read_file(str(path)))
13 label = path.parent.parent.name
14
15
16
17 # image = cv2.imread(str(path))
18 # # image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
19 # image = cv2.resize(image, (IMG_HEIGHT, IMG_WIDTH))
20 # image = image / 255
21 # # image = preprocess_input(image)
22 image = prep_external_image(str(path), preprocess_vgg)
23 net_input = utils.extend_tensor(image, 0)
24 prediction = model(net_input)
25 prediction = prediction[0].numpy()
26
27 valued_arr = []
28
29 for idx, name in enumerate(labels):
30     valued_arr.append(f'{name}\n({prediction[idx]:.2%})')
31
32 fig, ax = plt.subplots(
33     nrows=1, ncols=2, width_ratios=[0.4, 0.6], figsize=(20, 10)
34 )
35
36 # pltDisplay(image * 255, ax=ax[0]) # 1 channel
37 pltDisplay(image / 255, ax=ax[0]) # 3 channels
38
39 ax[1].bar(valued_arr, prediction)
40 plt.xticks(fontsize=15)
41 plt.yticks(fontsize=15)
42 plt.title(f'Correct label: {label}', fontsize=20)
43 # plt.xlabel('Predicted class')
44 # plt.ylabel('Percentage')
45 plt.show()

```

Generated/Frames/fearful/21/01-01-06-01-02-01-21-086-01.jpg



✓ Export the model with preprocessing

The model's not very easy to use if you have to apply those preprocessing steps before passing data to the model for inference. So build an end-to-end version: