


Import frames and train, test, evaluate network on facial images

✕ Importing Dependencies

```
1 from google.colab import drive
2 drive.mount('/content/drive', force_remount=True)
```

 Mounted at /content/drive

```
1 PROJECT_FOLDER = '/content/drive/MyDrive/CNN_Emotion_Classification/'
```

```
1 # %cd /content/drive/MyDrive/CNN_Emotion_Classification/
```

```
1 # copy content in main folder
2 ! cp -a {PROJECT_FOLDER}. ./
```

```
1 !pip install -r Requirements/pireqs_opencv_contrib_env.txt
```

 [Show hidden output](#)

Double-click (or enter) to edit

```
1 # !pip install keras_applications
```

```
1 from display import pltDisplay
2 from pathlib import Path
3 from matplotlib.colors import ListedColormap
4 from sklearn.metrics import classification_report, roc_curve, auc, roc_auc_score
5 from sklearn.preprocessing import LabelBinarizer
6 from sklearn.model_selection import train_test_split
7 # from sorting import human_sort
8 from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, Dense, Flatten
9 from tensorflow.keras.layers import BatchNormalization, Dropout, Activation, ReLU, Softmax, GlobalAveragePooling2D
10 from tensorflow.keras.models import Model, Sequential
11 from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau
12 from tensorflow.keras.regularizers import l2
13 # import IDython_display as id
```

```

13 # import IPython.display as ipd
14
15
16 import constants as const
17 import csv
18 import cv2
19 import json
20 import logging
21 import gc
22 import matplotlib.pyplot as plt
23 import numpy as np
24 import pandas as pd
25 import random
26 import seaborn as sns
27 import tensorflow as tf
28 import subprocess
29 import tensorflow.keras.backend as K
30 import time
31 import utils
32
33 from keras_vggface.vggface import VGGFace
34 from tensorflow.keras.applications import EfficientNetB0
35 from tensorflow.keras.applications.vgg16 import VGG16
36 from tensorflow.keras.applications.vgg16 import preprocess_input as preprocess_imagenet
37 from keras_vggface.utils import preprocess_input as preprocess_vggface

```

```
1 print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
```

 Num GPUs Available: 1

```

1 gpus = tf.config.list_physical_devices('GPU')
2 if gpus:
3     try:
4         # Currently, memory growth needs to be the same across GPUs
5         for gpu in gpus:
6             print("Name:", gpu.name, " Type:", gpu.device_type)
7             print(tf.config.experimental.get_device_details(gpu))
8             print(tf.config.experimental.get_memory_info('GPU:0'))
9             # tf.config.experimental.set_memory_growth(gpu, True)
10    logical_gpus = tf.config.list_logical_devices('GPU')
11    print(len(gpus), "Physical GPUs,", len(logical_gpus), "Logical GPUs")
12 except RuntimeError as e:
13     # Memory growth must be set before GPUs have been initialized
14     print(e)

```

```
➤ Name: /physical_device:GPU:0    Type: GPU
{'compute_capability': (7, 5), 'device_name': 'Tesla T4'}
{'current': 0, 'peak': 0}
1 Physical GPUs, 1 Logical GPUs
```

```
1 # unzip archive in folder
2 !7z x Generated/Frames_300.zip -oGenerated/
```

```
➤
7-Zip [64] 16.02 : Copyright (c) 1999-2016 Igor Pavlov : 2016-05-21
p7zip Version 16.02 (locale=en_US.UTF-8,Utf16=on,HugeFiles=on,64 bits,2 CPUs Intel(R) Xeon(R) CPU @ 2.00GHz (50653),ASM,AES-NI)

Scanning the drive for archives:
1 file, 4154172627 bytes (3962 MiB)

Extracting archive: Generated/Frames_300.zip
--
Path = Generated/Frames_300.zip
Type = zip
Physical Size = 4154172627
64-bit = +

Everything is Ok

Folders: 202
Files: 158556
Size:      4143561934
Compressed: 4154172627
```

```
1 utils.modules_info()
```

```
➤
OpenCV:
  Version: 4.8.0
Tensorflow:
  Version: 2.15.0
```

```
1 log_file = Path(const.logs_path, 'FACER.log')
2 logging.basicConfig(
3     format='%(asctime)s %(message)s',
4     filemode='a',
5     filename=log_file,
6     encoding='utf-8',
7     level=logging.INFO,
8     force=True
9 )
```

✓ Importing Dataset

```
1 data_df = pd.read_csv(Path(const.csv_path, 'dataset.csv'))
```

✓ Frame

✓ Preparing Data

✓ Dataset Creation for ML

```
1 IMG_WIDTH = IMG_HEIGHT = 224
2 IMG_CHANNELS = 3
3 SEED = 42
4 BATCH_SIZE = 128
5 VALIDATION_SPLIT = 0.2
6 EMOTIONS_LABELS = const.EMOTIONS_LABELS # RAVDESS emotion labels
```

```
1 EMOTIONS_LABELS
```

```
↔ ['neutral', 'calm', 'happy', 'sad', 'angry', 'fearful', 'disgust', 'surprised']
```

```
1 TOTAL_ELEMENTS = const.DATASET_TOTAL_ELEMENTS
2 label_names = const.EMOTIONS_LABELS_SORTED.copy()
3 # label_names.remove('neutral')
4 # label_names.remove('calm')
5 # label_names.remove('surprised')
```

```
1 label_names_gender = []
2 for em in label_names:
3     label_names_gender.append(em + '_female')
4     label_names_gender.append(em + '_male')
```

```
1 def get_gender(fname):
2
3     gender = int(utils.get_actor(fname)) % 2
4     gender_str = 'male' if (gender == 1) else 'female'
5     return [gender, gender_str]
```

▼ Dataset Creation - NEW

```
1 actors_labels = [f'{i:02d}' for i in range(1, 25)]
2 dist_idxes = {
3     '1': [slice(0, 16), slice(16, 20), slice(20, 24)],
4     '2': [slice(8, 24), slice(4, 8), slice(0, 4)],
5     '3': [slice(4, 20), slice(20, 24), slice(0, 4)]
6 }
```

```
1 # Split actors in train, validation, test
2 dist_n = 1
3 train_idxes, val_idxes, test_idxes = [actors_labels[i] for i in dist_idxes[str(dist_n)]]
4
5 print(train_idxes, val_idxes, test_idxes)
```

🔄 ['01', '02', '03', '04', '05', '06', '07', '08', '09', '10', '11', '12', '13', '14', '15', '16'] ['17', '18', '19', '20'] ['21', '22', '23', '24']

```

1 def make_dataset(path, actors_idx, talk_frame=False, acted_frame=False,
2                 undersampling=False, preprocess_vgg=True, shuffle=False,
3                 gender_classes=False, sampling=1):
4     # gender classes -> create a men/woman category for each emotion
5     def parse_image(filename):
6         image = tf.io.read_file(filename)
7         image = tf.image.decode_jpeg(image, channels=IMG_CHANNELS)
8         image = tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH])
9         if (preprocess_vgg == 'Imagenet'):
10             image = preprocess_imagenet(image)
11         elif (preprocess_vgg == 'VGGFace'):
12             # function does not accept tensor
13             # so convert to np and then to tensor
14             # image = np.array(image)
15             image = preprocess_vggface(image, version=2)
16             # image = tf.convert_to_tensor(image)
17         else:
18             image = image / 255
19
20         return image
21
22     # label_names = label_names if gender_classes else label_names_gender
23
24     filenames = []
25     talk_regex = '*-01.jpg' if talk_frame else '*.jpg'
26     acted_regex = '02' if acted_frame else '*'
27     gen_regex = f'*-*-*-{acted_regex}-*-*-*-{talk_regex}'
28
29     file_dict = dict()
30     file_dict_gender = dict()
31
32     for label in sorted(label_names):
33         file_dict[label] = []
34
35     for label in sorted(label_names_gender):
36         file_dict_gender[label] = []
37
38
39     for label in label_names:
40         for actor in actors_idx:
41             for file in Path(path, label, actor).glob(f'{gen_regex}'):
42                 gender = get_gender(str(file))[1]
43                 lab = label + f'_{gender}'
44                 file_dict_gender[lab].append(str(file))
45                 file_dict[label].append(str(file))
46

```



```

47     if shuffle:
48         for label, item in file_dict_gender.items():
49             logging.info(f'Label: {label}')
50             logging.info(f'Array len: {len(item)}')
51             random.Random(SEED).shuffle(item)
52
53
54     # NO GENDER
55
56
57     arr_len = [len(arr) for arr in file_dict.values()]
58
59     if undersampling:
60         filenames = [arr[:min(arr_len)] for arr in file_dict.values()]
61     else:
62         filenames = [arr for arr in file_dict.values()]
63
64     filenames = sum(filenames, [])
65
66     if shuffle:
67         random.Random(SEED).shuffle(filenames)
68
69     labels = []
70
71     if (gender_classes):
72
73         for elem in filenames:
74             cl = utils.get_class_string(str(elem))
75             gender = get_gender(str(elem))[1]
76             lab = cl + f'_{gender}'
77             labels.append(label_names_gender.index(lab))
78     else:
79         labels = [
80             label_names.index(EMOTIONS_LABELS[int(utils.get_class(elem)) - 1])
81             for elem in filenames
82         ]
83
84     if (sampling < 1):
85         filenames, _, labels, _ = train_test_split(
86             filenames, labels, train_size=sampling, random_state=SEED
87         )
88
89     filenames_ds = tf.data.Dataset.from_tensor_slices(filenames)
90     labels_ds = tf.data.Dataset.from_tensor_slices(labels)
91
92     images_ds = filenames_ds.map(

```



```

93         parse_image, num_parallel_calls=tf.data.experimental.AUTOTUNE
94     )
95     ds = tf.data.Dataset.zip((images_ds, labels_ds))
96
97     return [ds, filenames]

```

```

1 sampling_rate = 1
2 gender_classes = False
3 talk_frame = True
4 acted_frame = False
5 preprocess_vgg = 'Imagenet' # False, Imagenet or VGGFace
6 NUM_CLASSES = len(label_names_gender) if gender_classes else len(label_names)
7
8
9 train_ds, train_files = make_dataset(
10     const.frames_path, train_idx, talk_frame=talk_frame, acted_frame=acted_frame,
11     preprocess_vgg=preprocess_vgg, shuffle=True, gender_classes=gender_classes,
12     sampling=sampling_rate
13 )
14
15 val_ds, val_files = make_dataset(
16     const.frames_path, val_idx, talk_frame=talk_frame, acted_frame=acted_frame,
17     preprocess_vgg=preprocess_vgg, gender_classes=gender_classes,
18     sampling=sampling_rate
19 )
20
21 test_ds, test_files = make_dataset(
22     const.frames_path, test_idx, talk_frame=talk_frame, acted_frame=acted_frame,
23     preprocess_vgg=preprocess_vgg, gender_classes=gender_classes,
24     sampling=sampling_rate
25 )
26
27 labels = label_names if not gender_classes else label_names_gender

```

```

1 # train_ds = train_ds[:31882]
2 # train_files = train_files[:31882]
3
4 # val_ds = val_ds[:8426]
5 # val_files = val_files[:8426]
6
7 # test_ds = test_ds[:8301]
8 # test_files = test_files[:8301]

```

```
1 assert len(train_ds) == len(train_files), len(train_files)
2 assert len(val_ds) == len(val_files), len(val_files)
3 assert len(test_ds) == len(test_files), len(test_files)
```

```
1 train_ds_elements = len(train_ds)
2 test_ds_elements = len(test_ds)
3 val_ds_elements = len(val_ds)
```

```
1 print(f'train_ds samples: {train_ds_elements}')
2 print(f'test_ds samples: {test_ds_elements}')
3 print(f'val_ds samples: {val_ds_elements}')
```

```
➡ train_ds samples: 67906
   test_ds samples: 17994
   val_ds samples: 17642
```

✓ Build and train the model

Add operations to reduce read latency while training the model:

`ds.batch` Combines consecutive elements of the dataset into batches. The components of the resulting element will have an additional outer dimension, which will be *batch_size*

`ds.cache` Caches the elements in this dataset.

`ds.prefetch` Allows later elements to be prepared while the current element is being processed. This often improves latency and throughput, at the cost of using additional memory to store prefetched elements.

```
1 def configure_for_performance(ds, batch_size=BATCH_SIZE):
2     ds = ds.batch(batch_size)
3     # ds = ds.cache()
4     # ds = ds.shuffle(buffer_size=1000)
5     # ds = ds.repeat()
6     ds = ds.prefetch(buffer_size=tf.data.AUTOTUNE)
7     return ds
```

```
1 train_ds = configure_for_performance(train_ds)
2 val_ds = configure_for_performance(val_ds)
3 test_ds = configure_for_performance(test_ds)
```

```
1 for example_images, example_labels in train_ds.take(1):
2     print(example_images.shape)
3     print(example_labels.shape)
```


```
↔ (128, 224, 224, 3)
   (128,)
```

```
1 plt.figure(figsize=(16, 10))
2 rows = 2
3 cols = 2
4 n = rows * cols
5 for i in range(n):
6     plt.subplot(rows, cols, i + 1)
7     image = example_images[i]
8     plt.imshow(image) # 3 channels
9     # plt.imshow(image * 255, cmap='gray', vmin=0, vmax=255) # 1 channel
10    plt.title(f'{label_names[example_labels[i]]}\n{train_files[i]}')
```



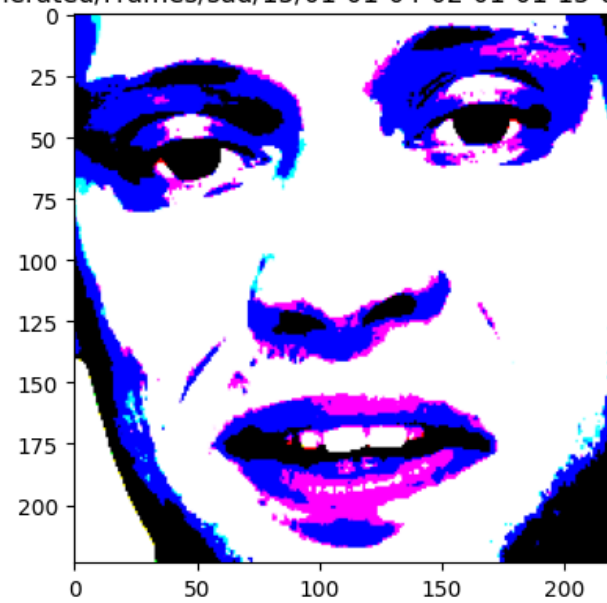
neutral

Generated/Frames/neutral/11/01-01-01-01-02-01-11-035-01.jpg




sad

Generated/Frames/sad/15/01-01-04-02-01-01-15-068-01.jpg



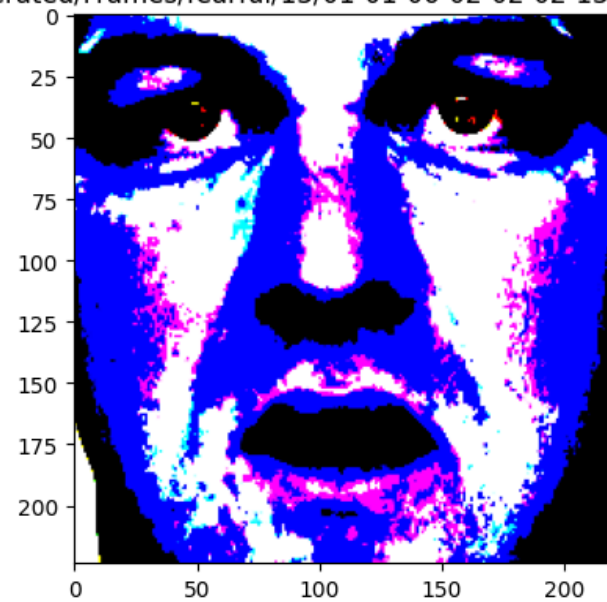
disgust


Generated/Frames/disgust/04/01-01-07-01-01-01-04-071-01.jpg



fearful

Generated/Frames/fearful/13/01-01-06-02-02-02-13-031-01.jpg





```

1 def create_Bilotti_CNN(name='Bilotti_CNN'):
2
3     inputs = Input(shape=(IMG_HEIGHT, IMG_WIDTH, IMG_CHANNELS))
4
5     conv1 = Conv2D(32, kernel_size=(3, 3), activation='relu')(inputs)
6     conv2 = Conv2D(32, kernel_size=(3, 3), activation='relu')(conv1)
7     pool1 = MaxPooling2D(pool_size=(2, 2))(conv2)
8
9     conv3 = Conv2D(64, kernel_size=(3, 3), activation='relu')(pool1)
10    conv4 = Conv2D(64, kernel_size=(3, 3), activation='relu')(conv3)
11    pool2 = MaxPooling2D(pool_size=(2, 2))(conv4)
12
13    conv4 = Conv2D(18, kernel_size=(3, 3), activation='relu')(pool2)
14    conv5 = Conv2D(18, kernel_size=(3, 3), activation='relu')(conv4)
15    conv6 = Conv2D(18, kernel_size=(3, 3), activation='relu')(conv5)
16    pool3 = MaxPooling2D(pool_size=(2, 2))(conv6)
17
18    conv7 = Conv2D(56, kernel_size=(3, 3), activation='relu')(pool3)
19    conv8 = Conv2D(56, kernel_size=(3, 3), activation='relu')(conv7)
20    conv9 = Conv2D(56, kernel_size=(3, 3), activation='relu')(conv8)
21    pool4 = MaxPooling2D(pool_size=(2, 2))(conv9)
22
23    conv10 = Conv2D(51, kernel_size=(3, 3), activation='relu')(pool4)
24    conv11 = Conv2D(51, kernel_size=(3, 3), activation='relu')(conv10)
25    conv12 = Conv2D(51, kernel_size=(3, 3), activation='relu')(conv11)
26    pool5 = MaxPooling2D(pool_size=(2, 2))(conv12)
27
28    flatten = Flatten()(pool5)
29
30    dense1 = Dense(2048, activation='relu')(flatten)
31    drop1 = Dropout(0.25)(dense1)
32
33    dense2 = Dense(1024, activation='relu')(drop1)
34    drop2 = Dropout(0.4)(dense2)
35
36    output = Dense(NUM_CLASSES, activation='softmax')(drop2)
37
38    model = Model(inputs, output)
39
40    model._name = name
41
42    return model

```

```
1 def create_VGG16_Imagenet(name='VGG16_Imagenet'):
2
3     base_model = VGG16(
4         weights='imagenet',
5         include_top=False,
6         input_shape=(IMG_WIDTH, IMG_HEIGHT, IMG_CHANNELS)
7     )
8     base_model.trainable = False # Not trainable weights
9
10    flatten_layer = Flatten()
11    dense_layer_1 = Dense(2048, activation='relu')
12    drop_1 = Dropout(0.4)
13    dense_layer_2 = Dense(1024, activation='relu')
14    drop_2 = Dropout(0.4)
15    dense_layer_3 = Dense(512, activation='relu')
16    drop_3 = Dropout(0.4)
17    prediction_layer = Dense(NUM_CLASSES, activation='softmax')
18
19    model = Sequential([
20        base_model,
21        flatten_layer,
22        dense_layer_1,
23        drop_1,
24        dense_layer_2,
25        drop_2,
26        dense_layer_3,
27        drop_3,
28        prediction_layer
29    ])
30
31    model._name = name
32
33    return model
```

```
1 def create_EfficientNetB0_Imagenet(name='EfficientNetB0_Imagenet'):  
2  
3     nb_class = NUM_CLASSES  
4  
5     inputs = Input(shape=(IMG_WIDTH, IMG_HEIGHT, IMG_CHANNELS))  
6     model = EfficientNetB0(  
7         include_top=False, input_tensor=inputs, weights="imagenet"  
8     )  
9  
10    # Freeze the pretrained weights  
11    model.trainable = False  
12  
13    # Rebuild top  
14    x = GlobalAveragePooling2D(name="avg_pool")(model.output)  
15    x = BatchNormalization()(x)  
16  
17    top_dropout_rate = 0.2  
18    x = Dropout(top_dropout_rate, name="top_dropout")(x)  
19    outputs = Dense(nb_class, activation="softmax", name="pred")(x)  
20  
21    model._name = name  
22    # Compile  
23    model = Model(inputs, outputs, name="EfficientNet")  
24  
25    return model
```



```

1 def create_VGG16_VGGFACE(name='VGG16_VGGFACE'):
2     nb_class = NUM_CLASSES
3
4     vgg_model = VGGFace(
5         include_top=False, weights='vggface', input_shape=(IMG_WIDTH, IMG_HEIGHT, IMG_CHANNELS)
6     )
7     last_layer = vgg_model.get_layer('pool5').output
8     x = Flatten(name='flatten')(last_layer)
9
10    x = Dense(512, activation='relu', name='fc6')(x)
11    x = Dropout(0.35)(x)
12    x = Dense(256, activation='relu', name='fc7')(x)
13    x = Dropout(0.35)(x)
14    x = Dense(128, activation='relu', name='fc8')(x)
15    x = Dropout(0.35)(x)
16
17    out = Dense(nb_class, activation='softmax', name='fc9')(x)
18
19    custom_vgg_model = Model(vgg_model.input, out)
20    custom_vgg_model._name = name
21
22    return custom_vgg_model

```

```

1 tf.keras.backend.clear_session() # clear all precedent models and sessions

```

```

1 check_path = 'checkpoint.weights.h5'
2 checkpointer = ModelCheckpoint(
3     check_path, monitor='val_accuracy', verbose=1, save_best_only=True,
4     save_weights_only=False, mode='auto', save_freq='epoch'
5 )

```

```

1 # model = create_cnn_model()
2 # model = medium_model()
3 # model = create_VGG16_Imagenet()
4 # model = create_VGG16_VGGFACE()
5 model = create_EfficientNetB0_Imagenet()
6 # model = create_grigorasi_model()
7 # model = create_Bilotti_CNN()
8 model.summary()

```




```
block6b_se_excite (Multipl (None, 7, 7, 1152)      0      ['block6b_activation[0][0]',  
y)                                                  'block6b_se_expand[0][0]']
```

```
1 # Define training callbacks  
2  
3 class TimeHistory(tf.keras.callbacks.Callback):  
4     def on_train_begin(self, logs={}):  
5         self.times = []  
6  
7     def on_epoch_begin(self, batch, logs={}):  
8         self.epoch_time_start = time.time()  
9  
10    def on_epoch_end(self, batch, logs={}):  
11        self.times.append(time.time() - self.epoch_time_start)  
12  
13  
14 early_stopping_callback = tf.keras.callbacks.EarlyStopping(  
15     verbose=1,  
16     patience=5,  
17     restore_best_weights=True  
18 )  
19  
20 reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, verbose=1,  
21                               patience=3, min_lr=0)
```

```
1 METRICS = ['accuracy']
```

```
1 model.compile(  
2     optimizer=tf.keras.optimizers.Adam(  
3         learning_rate=1e-2  
4     ),  
5     # optimizer=tf.keras.optimizers.SGD(), # for VGG16_VGGFACE  
6     loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),  
7     metrics=METRICS,  
8 )
```

✓ Train Model

```

1 EPOCHS = 100
2 time_callback = TimeHistory()
3 history = model.fit(
4     x=train_ds,
5     validation_data=val_ds,
6     epochs=EPOCHS,
7     callbacks=[time_callback, early_stopping_callback, reduce_lr]
8 )

```



```

Epoch 1/100
531/531 [=====] - 367s 664ms/step - loss: 1.1273 - accuracy: 0.6550 - val_loss: 2.0916 - val_accuracy: 0.4065 - lr: 0.0100
Epoch 2/100
531/531 [=====] - 323s 609ms/step - loss: 0.9074 - accuracy: 0.7043 - val_loss: 2.0951 - val_accuracy: 0.4336 - lr: 0.0100
Epoch 3/100
531/531 [=====] - 330s 622ms/step - loss: 0.8689 - accuracy: 0.7105 - val_loss: 2.1393 - val_accuracy: 0.4371 - lr: 0.0100
Epoch 4/100
531/531 [=====] - ETA: 0s - loss: 0.8555 - accuracy: 0.7128
Epoch 4: ReduceLROnPlateau reducing learning rate to 0.0019999999552965165.
531/531 [=====] - 316s 595ms/step - loss: 0.8555 - accuracy: 0.7128 - val_loss: 2.0959 - val_accuracy: 0.4111 - lr: 0.0100
Epoch 5/100
531/531 [=====] - 318s 599ms/step - loss: 0.7495 - accuracy: 0.7473 - val_loss: 1.8994 - val_accuracy: 0.4540 - lr: 0.0020
Epoch 6/100
531/531 [=====] - 322s 606ms/step - loss: 0.7267 - accuracy: 0.7526 - val_loss: 1.8757 - val_accuracy: 0.4581 - lr: 0.0020
Epoch 7/100
531/531 [=====] - 317s 598ms/step - loss: 0.7259 - accuracy: 0.7546 - val_loss: 1.8582 - val_accuracy: 0.4761 - lr: 0.0020
Epoch 8/100
531/531 [=====] - 321s 604ms/step - loss: 0.7227 - accuracy: 0.7562 - val_loss: 1.8993 - val_accuracy: 0.4639 - lr: 0.0020
Epoch 9/100
531/531 [=====] - 337s 635ms/step - loss: 0.7229 - accuracy: 0.7574 - val_loss: 1.8278 - val_accuracy: 0.4675 - lr: 0.0020
Epoch 10/100
531/531 [=====] - 320s 603ms/step - loss: 0.7174 - accuracy: 0.7587 - val_loss: 1.8641 - val_accuracy: 0.4719 - lr: 0.0020
Epoch 11/100
531/531 [=====] - 332s 626ms/step - loss: 0.7194 - accuracy: 0.7569 - val_loss: 1.9551 - val_accuracy: 0.4593 - lr: 0.0020
Epoch 12/100
531/531 [=====] - ETA: 0s - loss: 0.7180 - accuracy: 0.7581
Epoch 12: ReduceLROnPlateau reducing learning rate to 0.0003999999724328518.
531/531 [=====] - 318s 599ms/step - loss: 0.7180 - accuracy: 0.7581 - val_loss: 2.0049 - val_accuracy: 0.4496 - lr: 0.0020
Epoch 13/100
531/531 [=====] - 330s 621ms/step - loss: 0.6904 - accuracy: 0.7662 - val_loss: 1.9980 - val_accuracy: 0.4525 - lr: 4.0000e-04
Epoch 14/100
531/531 [=====] - ETA: 0s - loss: 0.6872 - accuracy: 0.7678Restoring model weights from the end of the best epoch: 9.
531/531 [=====] - 319s 601ms/step - loss: 0.6872 - accuracy: 0.7678 - val_loss: 1.9956 - val_accuracy: 0.4503 - lr: 4.0000e-04
Epoch 14: early stopping

```

```

1 EPOCHS = len(time_callback.times)

```


```
1 # Create model path
2 model_path = Path(const.models_path, 'Frame', model._name)
3 run_folders = list(Path(const.models_path, 'Frame', model._name).glob('Run_*'))
4
5 if not run_folders:
6     model_path = Path(model_path, 'Run_1')
7 else:
8     last_run = run_folders.pop()
9     last_run_idx = Path(last_run).name.split('_')[-1]
10    model_path = Path(model_path, f'Run_{int(last_run_idx) + 1}')
11
12 model_path.mkdir(parents=True, exist_ok=False)
```

```
1 # Save info on the indexes used for train, val and test
2 ds_info_path = Path(model_path, f'{model._name}_dataset.txt')
3 with open(ds_info_path, 'w+', newline='') as res_file:
4     res_file.write(f'Train indexes: {train_idxs}\n')
5     res_file.write(f'Train files: {train_ds_elements}\n')
6     res_file.write(f'Val indexes: {val_idxs}\n')
7     res_file.write(f'Val files: {val_ds_elements}\n')
8     res_file.write(f'Test indexes: {test_idxs}\n')
9     res_file.write(f'Test files: {test_ds_elements}\n')
```


```
1 metrics = history.history
```

```
1 mod_loss = metrics['loss']
2 mod_val_loss = metrics['val_loss']
3 mod_accuracy = metrics['accuracy']
4 mod_val_accuracy = metrics['val_accuracy']
5 # mod_f1 = metrics['fBeta_score']
6 # mod_val_f1 = metrics['val_fBeta_score']
7
8 mod_mean_loss = np.mean(mod_loss)
9 mod_mean_val_loss = np.mean(mod_val_loss)
10 mod_mean_accuracy = np.mean(mod_accuracy)
11 mod_mean_val_accuracy = np.mean(mod_val_accuracy)
12 # mod_mean_f1 = np.mean(mod_f1)
13 # mod_mean_val_f1 = np.mean(mod_val_f1)
```

```
1 # Save Loss
2 plt.title('Model loss')
3 plt.plot(history.epoch, mod_loss, mod_val_loss)
4 plt.legend(['loss', 'val_loss'])
5 plt.ylim([0, max(plt.ylim())])
6 plt.xlabel('Epoch')
7 plt.ylabel('Loss [CrossEntropy]')
8 plt.savefig(Path(model_path, 'loss.png'))
```



```
1 # Save Accuracy
2 plt.title('Model accuracy')
3 plt.plot(
4     history.epoch,
5     100 * np.array(mod_accuracy),
6     100 * np.array(mod_val_accuracy)
7 )
8 plt.legend(['accuracy', 'val_accuracy'])
9 plt.ylim([0, 100])
10 plt.xlabel('Epoch')
11 plt.ylabel('Accuracy [%]')
12 plt.savefig(Path(model_path, 'accuracy.png'))
```



Unsupported Cell Type. Double-Click to inspect/edit the content.

Unsupported Cell Type. Double-Click to inspect/edit the content.

✓ Evaluate Model

```
1 model_eval = model.evaluate(test_ds, return_dict=True)
```

↔ 141/141 [=====] - 71s 497ms/step - loss: 1.8451 - accuracy: 0.4229

```
1 # Save model
2 model.save(Path(model_path, f'{model._name}.keras'), overwrite=False)
3 # Save history
4 np.save(Path(model_path, f'{model._name}_history.npy'), history)
5 # Save model image
6 model_img = tf.keras.utils.plot_model(
7     model, Path(model_path, f'{model._name}.png'), show_shapes=True,
8     show_layer_names=True, show_layer_activations=True
9 )
```

```
1 model_eval
```

↔ {'loss': 1.8450863361358643, 'accuracy': 0.4228631854057312}

```
1 test_loss = model_eval['loss']
2 test_accuracy = model_eval['accuracy']
3 # test_f1 = model_eval['fBeta_score']
4 mean_epoch_time = np.mean(time_callback.times)
```

```
1 # Salvataggio informazioni modello
2 model_save_path = Path(model_path, f'{model._name}_result.txt')
3 with open(model_save_path, 'w+', newline='') as res_file:
4     res_file.write(f'BATCH: {BATCH_SIZE}\n')
5     res_file.write(f'Train loss: {str(mod_loss)}\n')
6     res_file.write(f'val_loss: {str(mod_val_loss)}\n')
7     res_file.write(f'Train accuracy: {str(mod_accuracy)}\n')
8     res_file.write(f'Train val_accuracy: {str(mod_val_accuracy)}\n')
9     # res_file.write(f'Train f1_score: {str(mod_f1)}\n')
10    # res_file.write(f'Train val_f1_score: {str(mod_val_f1)}\n')
11    res_file.write(f'Test loss: {str(test_loss)}\n')
12    res_file.write(f'Test accuracy: {str(test_accuracy)}\n')
13    # res_file.write(f'Test f1_score: {str(test_f1)}\n')
14    res_file.write(f'Mean epoch time: {str(mean_epoch_time)}')
```

```
1 # Salvataggio informazioni generali modelli
```



```

2 with open(Path(const.models_path, 'Frame', 'models.csv'), 'a+') as csvfile:
3     filewriter = csv.writer(
4         csvfile, delimiter=';', quotechar='|', quoting=csv.QUOTE_MINIMAL
5     )
6
7     # filewriter.writerow(
8     #     ["Model Name", "Epochs", "% Validation", "% Test set",
9     #     "Train loss", "Train accuracy", "Val loss", "Val accuracy",
10    #     "Test loss", "Test accuracy", "Mean epoch time", "Note"]
11    # )
12    test_ds_perc = utils.trunc((test_ds_elements * 100) / TOTAL_ELEMENTS, 2)
13    val_ds_perc = utils.trunc((val_ds_elements * 100) / TOTAL_ELEMENTS, 2)
14    full_path = str(Path(model._name, model_path.name))
15    filewriter.writerow(
16        [full_path, EPOCHS, val_ds_perc, test_ds_perc,
17         mod_loss, mod_accuracy, mod_val_loss, mod_val_accuracy,
18         test_loss, test_accuracy, mean_epoch_time, '']
19    )

```

```

1 for test_images, test_labels in test_ds.take(1):
2     print(test_images.shape)
3     print(test_labels.shape)

```

```

➡ (128, 224, 224, 3)
  (128,)

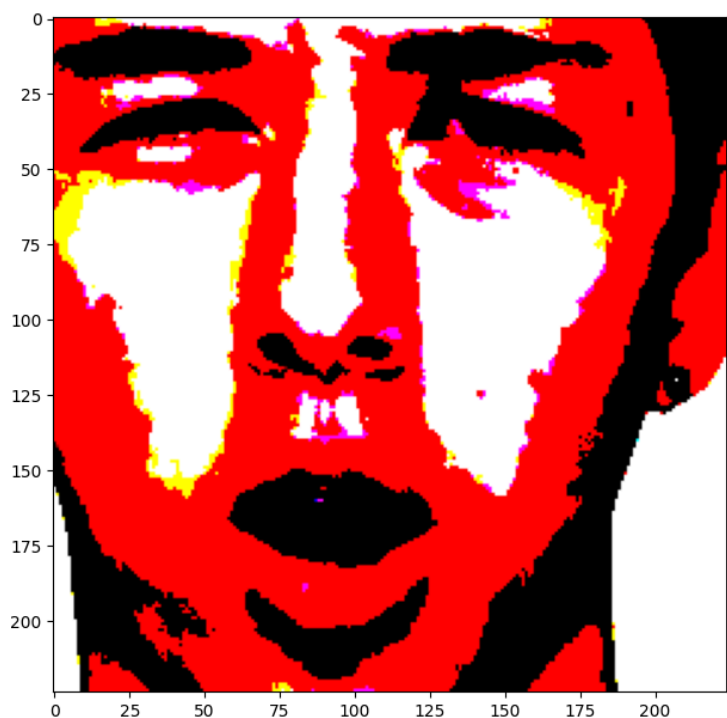
```

```


1 gen = np.random.default_rng(seed=None)
2 idx = gen.integers(0, len(test_images))
3 print(test_files[idx])
4
5 image = test_images[idx]
6 label = test_labels[idx]
7
8 net_input = utils.extend_tensor(image, 0)
9 prediction = model(net_input)
10 prediction = prediction[0].numpy()
11
12 valued_arr = []
13
14 for idx, name in enumerate(label_names):
15     valued_arr.append(f'{name}\n({prediction[idx]:.2%})')
16
17 fig, ax = plt.subplots(
18     nrows=1, ncols=2, width_ratios=[0.4, 0.6], figsize=(20, 10)
19 )

```

```
20
21 pltDisplay(image * 255, ax=ax[0]) # 1 channel
22 # pltDisplay(image, ax=ax[0])
23
24 ax[1].bar(valued_arr, prediction)
25 plt.xticks(fontsize=15)
26 plt.yticks(fontsize=15)
27 plt.title(f'Correct label: {label_names[label]}', fontsize=20)
28 # plt.xlabel('Predicted class')
29 # plt.ylabel('Percentage')
30 plt.show()
```



Correct label: angry



Display a confusion matrix

Use a [confusion matrix](#) to check how well the model did classifying each of the commands in the test set:

Unsupported Cell Type. Double-Click to inspect/edit the content.

```
1 y_pred = model.predict(test_ds)
```

🔄 141/141 [=====] - 76s 510ms/step


```
1 y_pred = tf.argmax(y_pred, axis=1, output_type=tf.int32)
```

```
1 y_true = tf.concat(list(test_ds.map(lambda _, lab: lab)), axis=0)
```

```
1 report = classification_report(  
2     y_true, y_pred, target_names=label_names,  
3     output_dict=True, zero_division='warn'  
4 )
```


```
1 rep_to_csv = pd.DataFrame(data=report).transpose()
```

```
1 fig, ax = plt.subplots(figsize=(10, 10))  
2 ax.xaxis.tick_top()  
3 sns.heatmap(rep_to_csv.iloc[:NUM_CLASSES, :3],  
4             cbar=True,  
5             square=False,  
6             annot=True,  
7             annot_kws={'size': 15},  
8             fmt='.2g',  
9             linewidths=0.5)  
10 plt.savefig(Path(model_path, f'{model._name}_f1_score.png'))  
11 plt.show()  
12  
13  
14 with sns.axes_style('white'):  
15     fig, ax = plt.subplots(figsize=(3, 5))  
16     ax.xaxis.tick_top()  
17     sns.heatmap(rep_to_csv.iloc[:NUM_CLASSES, 3:],  
18               cbar=False,  
19               square=False,  
20               annot=True,  
21               annot_kws={'size': 15},  
22               fmt='.4g',  
23               cmap=ListedColormap(['purple']),  
24               linewidths=0.5)  
25     plt.savefig(Path(model_path, f'{model._name}_support.png'))  
26     plt.show()
```



support





```
1 report_save_path = Path(model_path, f'{model._name}_report.csv')
2 rep_to_csv.to_csv(report_save_path)
```

Unsupported Cell Type. Double-Click to inspect/edit the content.

```
1 np.mean([report[c]['f1-score'] for c in list(report)[:NUM_CLASSES]])
```


0.4066562185019181

```
1 # True Test Distribution
2 # unique, counts = np.unique(y_true, return_counts=True)
3 # collections.Counter(y_true)
4 counts = [np.count_nonzero(y_true == idx) for idx in range(len(label_names))]
5 valued_arr = []
6 # for i in range(len(label_names)):
7 for idx, name in enumerate(label_names):
8     count = counts[idx]
9     valued_arr.append(f'{name}\n({count})')
10
11 fig = plt.subplots(figsize=(20, 10))
12 plt.bar(valued_arr, counts)
13 plt.xticks(fontsize=15)
```

```
14 plt.yticks(fontsize=15)
15 plt.title('Test True Distribution', fontsize=20)
16 # plt.xlabel('Predicted class')
17 # plt.ylabel('Percentage')
18 plt.savefig(Path(model_path, f'{model._name}_trueDist.png'))
19 plt.show()
```




Test True Distribution




```
1 # Predicted Test Distribution
2 # unique, counts = np.unique(y_pred, return_counts=True)
3 counts = [np.count_nonzero(y_pred == idx) for idx in range(len(label_names))]
4 valued_arr = []
5 unique_idx = 0
6 for idx, name in enumerate(label_names):
7     count = counts[idx]
8     valued_arr.append(f'{name}\n({count})')
9
10 fig = plt.subplots(figsize=(20, 10))
11 plt.bar(valued_arr, counts)
12 plt.xticks(fontsize=15)
13 plt.yticks(fontsize=15)
14 plt.title('Test result distribution', fontsize=20)
15 # plt.xlabel('Predicted class')
16 # plt.ylabel('Percentage')
17 plt.savefig(Path(model_path, f'{model._name}_predDist.png'))
18 plt.show()
```




Test result distribution



```
1 confusion_mtx = tf.math.confusion_matrix(y_true, y_pred)
2 fig, ax = plt.subplots(figsize=(10, 10))
3 sns.heatmap(confusion_mtx,
4             xticklabels=label_names,
5             yticklabels=label_names,
6             annot=True, fmt='.4g')
7 plt.xlabel('Prediction')
8 plt.ylabel('True Label')
9 plt.savefig(Path(model_path, f'{model._name}_heat.png'))
10 plt.show()
```



```

1 def multiclass_roc_auc_score(target, y_test, y_pred, average="macro"):
2     # function for scoring roc auc score for multi-class
3     lb = LabelBinarizer()
4     lb.fit(y_test)
5     y_test = lb.transform(y_test)
6     y_pred = lb.transform(y_pred)
7
8     if len(target) > 2:
9         for (idx, c_label) in enumerate(target):
10
11             fpr, tpr, thresholds = roc_curve(
12                 y_test[:, idx].astype(int),
13                 y_pred[:, idx]
14             )
15             c_ax.plot(
16                 fpr, tpr, label='%s (AUC:%0.2f)' % (c_label, auc(fpr, tpr))
17             )
18         else:
19             fpr, tpr, thresholds = roc_curve(
20                 y_test,
21                 y_pred
22             )
23             c_ax.plot(
24                 fpr, tpr, label='Model (AUC:%0.2f)' % (auc(fpr, tpr)), color='#ff7f0e'
25             )
26     c_ax.plot(fpr, fpr, color='b', linestyle='--', label='Random Guessing')
27
28     return roc_auc_score(y_test, y_pred, average=average)

```

```

1 # set plot figure size
2 fig, c_ax = plt.subplots(1, 1, figsize=(12, 8))
3
4 print('ROC AUC score:', multiclass_roc_auc_score(
5     label_names,
6     tf.reshape(y_true, (y_true.shape[0], 1)),
7     tf.reshape(y_pred, (y_pred.shape[0], 1))
8 ))
9
10 c_ax.legend()
11 c_ax.set_xlabel('False Positive Rate')
12 c_ax.set_ylabel('True Positive Rate')
13 plt.savefig(Path(model_path, f'{model._name}_ROC.png'))
14 plt.show()

```



ROC AUC score: 0.6686665117049038

