

Import frames and train, test, evaluate network on facial images

Importing Dependencies

```
1 from google.colab import drive
2 drive.mount('/content/drive', force_remount=True)
```

Mounted at /content/drive

```
1 PROJECT_FOLDER = '/content/drive/MyDrive/CNN_Emotion_Classification/'
```

```
1 # %cd /content/drive/MyDrive/CNN_Emotion_Classification/
```

```
1 # copy content in main folder
2 ! cp -a {PROJECT_FOLDER}. ./
```

```
1 !pip install -r Requirements/pireqs_opencv_contrib_env.txt
```

Requirement already satisfied: pooch>=1.1 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_contrib_env.tx

Requirement already satisfied: soxr>=0.3.2 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_contrib_env.t

Requirement already satisfied: typing-extensions>=4.1.1 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_

Requirement already satisfied: lazy-loader>=0.1 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_contrib_

Requirement already satisfied: msgpack>=1.0 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_contrib_env.

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_contr

Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_contrib_e

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_cont

Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_cont

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_contri

Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_contr

Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_c

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->-r Requirements/pireqs_opencv_contrib_env.t

Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->-r Requirements/pireqs_opencv_contrib_env

Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit_learn->-r Requirements/pireqs_opencv

Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib

Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_cont

Requirement already satisfied: flatbuffers>=23.5.26 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_c

Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/

Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_co

Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contr

Requirement already satisfied: ml-dtypes~0.2.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contr

Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contr

Requirement already satisfied: protobuf!=4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<5.0.0dev,>=3.20.3 in /usr/local/lib/python3.10/dis

Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contr

Requirement already satisfied: wrapt<1.15,>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_co

Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements

Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_co

Requirement already satisfied: tensorboard<2.16,>=2.15 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_openc

Requirement already satisfied: tensorflow-estimator<2.16,>=2.15.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/p

Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from astunparse>=1.6.0->tensorflow->-r Requirement

Requirement already satisfied: parso<0.9.0,>=0.8.3 in /usr/local/lib/python3.10/dist-packages (from jedi>=0.16->ipython->-r Requirements/pireqs_

Requirement already satisfied: llvmlite<0.42,>=0.41.0dev0 in /usr/local/lib/python3.10/dist-packages (from numba>=0.51.0->librosa->-r Requiremen

Requirement already satisfied: pyprocess>=0.5 in /usr/local/lib/python3.10/dist-packages (from pexpect>4.3->ipython->-r Requirements/pireqs_ope

Requirement already satisfied: platformdirs>=2.5.0 in /usr/local/lib/python3.10/dist-packages (from pooch>=1.1->librosa->-r Requirements/pireqs_

Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.10/dist-packages (from pooch>=1.1->librosa->-r Requirements/pireqs_ope

Requirement already satisfied: wcwidth in /usr/local/lib/python3.10/dist-packages (from prompt-toolkit!=3.0.0,!3.0.1,<3.1.0,>=2.0.0->ipython->-

Requirement already satisfied: cffi>=1.0 in /usr/local/lib/python3.10/dist-packages (from soundfile>=0.12.1->librosa->-r Requirements/pireqs_ope

Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.16,>=2.15->tensorflow->-r Re

Requirement already satisfied: google-auth-oauthlib<2,>=0.5 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.16,>=2.15->tensorflow

Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.16,>=2.15->tensorflow->-r Requirem

Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.16,>=2.15->t

Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.16,>=2.15->tensorflow->-r Requirem

Requirement already satisfied: pycparser in /usr/local/lib/python3.10/dist-packages (from cffi>=1.0->soundfile>=0.12.1->librosa->-r Requirements

Requirement already satisfied: cachetools<6.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3->tensorboard<2.16,>

Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3->tensorboard<2.16,>=

Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3->tensorboard<2.16,>=2.15->te

Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from google-auth-oauthlib<2,>=0.5->tensorboa

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->pooch>=1.1->librosa->

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->pooch>=1.1->librosa->-r Requireme

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->pooch>=1.1->librosa->-r Req

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->pooch>=1.1->librosa->-r Req

Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.10/dist-packages (from werkzeug>=1.0.1->tensorboard<2.16,>=2.15->tens

Requirement already satisfied: pyasn1<0.7.0,>=0.4.6 in /usr/local/lib/python3.10/dist-packages (from pyasn1-modules>=0.2.1->google-auth<3,>=1.6.

Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.10/dist-packages (from requests-oauthlib>=0.7.0->google-auth-oauthlib<2

Installing collected packages: jedi, keras_vggface, keras_applications

Successfully installed jedi-0.19.1 keras_applications-1.0.8 keras_vggface-0.6

Double-click (or enter) to edit

```
1
```

```
1 # !pip install keras_applications
```

```

1 from display import pltDisplay
2 from pathlib import Path
3 from matplotlib.colors import ListedColormap
4 from sklearn.metrics import classification_report, roc_curve, auc, roc_auc_score
5 from sklearn.preprocessing import LabelBinarizer
6 from sklearn.model_selection import train_test_split
7 # from sorting import human_sort
8 from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, Dense, Flatten
9 from tensorflow.keras.layers import BatchNormalization, Dropout, Activation, ReLU, Softmax
10 from tensorflow.keras.models import Model, Sequential
11 from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau
12 from tensorflow.keras.regularizers import l2
13 # import IPython.display as ipd
14
15
16 import constants as const
17 import csv
18 import cv2
19 import json
20 import logging
21 import gc
22 import matplotlib.pyplot as plt
23 import numpy as np
24 import pandas as pd
25 import random
26 import seaborn as sns
27 import tensorflow as tf
28 import subprocess
29 import tensorflow.keras.backend as K
30 import time
31 import utils
32
33 from keras_vggface.vggface import VGGFace
34 from tensorflow.keras.applications.vgg16 import VGG16
35 from tensorflow.keras.applications.vgg16 import preprocess_input as preprocess_imagenet
36 from keras_vggface.utils import preprocess_input as preprocess_vggface

```

```
1 print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
```

```
➦ Num GPUs Available: 1
```

```

1 gpus = tf.config.list_physical_devices('GPU')
2 if gpus:
3     try:
4         # Currently, memory growth needs to be the same across GPUs
5         for gpu in gpus:
6             print("Name:", gpu.name, " Type:", gpu.device_type)
7             print(tf.config.experimental.get_device_details(gpu))
8             print(tf.config.experimental.get_memory_info('GPU:0'))
9             # tf.config.experimental.set_memory_growth(gpu, True)
10        logical_gpus = tf.config.list_logical_devices('GPU')
11        print(len(gpus), "Physical GPUs,", len(logical_gpus), "Logical GPUs")
12    except RuntimeError as e:
13        # Memory growth must be set before GPUs have been initialized
14        print(e)

```

```

➦ Name: /physical_device:GPU:0   Type: GPU
{'compute_capability': (7, 5), 'device_name': 'Tesla T4'}
{'current': 0, 'peak': 0}
1 Physical GPUs, 1 Logical GPUs

```

```

1 # unzip archive in folder
2 !7z x Generated/Frames_300.zip -oGenerated/

```

```

➦
7-Zip [64] 16.02 : Copyright (c) 1999-2016 Igor Pavlov : 2016-05-21
p7zip Version 16.02 (locale=en_US.UTF-8,Utf16=on,HugeFiles=on,64 bits,2 CPUs Intel(R) Xeon(R) CPU @ 2.30GHz (306F0),ASM,AES-NI)

Scanning the drive for archives:
1 file, 4154172627 bytes (3962 MiB)

Extracting archive: Generated/Frames_300.zip
--
Path = Generated/Frames_300.zip
Type = zip
Physical Size = 4154172627
64-bit = +

Everything is Ok

Folders: 202
Files: 158556
Size:      4143561934
Compressed: 4154172627

```

```
1 utils.modules_info()
```



```
OpenCV:
  Version: 4.8.0
Tensorflow:
  Version: 2.15.0
```

```
1 log_file = Path(const.logs_path, 'FACER.log')
2 logging.basicConfig(
3     format='%(asctime)s %(message)s',
4     filemode='a',
5     filename=log_file,
6     encoding='utf-8',
7     level=logging.INFO,
8     force=True
9 )
```

✚ Importing Dataset

```
1 data_df = pd.read_csv(Path(const.csv_path, 'dataset.csv'))
```

✚ Frame

✚ Preparing Data

✚ Dataset Creation for ML

```
1 IMG_WIDTH = IMG_HEIGHT = 224
2 IMG_CHANNELS = 3
3 SEED = 42
4 BATCH_SIZE = 128
5 VALIDATION_SPLIT = 0.2
6 EMOTIONS_LABELS = const.EMOTIONS_LABELS # RAVDESS emotion labels
```

```
1 EMOTIONS_LABELS
```



```
['neutral', 'calm', 'happy', 'sad', 'angry', 'fearful', 'disgust', 'surprised']
```

```
1 TOTAL_ELEMENTS = const.DATASET_TOTAL_ELEMENTS
2 label_names = const.EMOTIONS_LABELS_SORTED.copy()
3 # label_names.remove('neutral')
4 # label_names.remove('calm')
5 # label_names.remove('surprised')
6
7 NUM_CLASSES = len(label_names)
```

✚ Dataset Creation - NEW

```
1 actors_labels = [f'{i:02d}' for i in range(1, 25)]
2 dist_idx = {
3     '1': [slice(0, 16), slice(16, 20), slice(20, 24)],
4     '2': [slice(8, 24), slice(4, 8), slice(0, 4)],
5     '3': [slice(4, 20), slice(20, 24), slice(0, 4)]
6 }
```

```
1 # Split actors in train, validation, test
2 dist_n = 1
3 train_idx, val_idx, test_idx = [actors_labels[i] for i in dist_idx[str(dist_n)]]
4
5 print(train_idx, val_idx, test_idx)
```



```
['01', '02', '03', '04', '05', '06', '07', '08', '09', '10', '11', '12', '13', '14', '15', '16'] ['17', '18', '19', '20'] ['21', '22', '23', '24']
```



```
1 def make_dataset(path, actors_idx, talk_frame=False, acted_frame=False,
2                 undersampling=False, preprocess_vgg=True, shuffle=False,
3                 sampling=1):
4
5     def parse_image(filename):
6         .....image = tf.io.read_file(filename)
7         .....image = tf.image.decode_jpeg(image, channels=IMG_CHANNELS)
8         .....image = tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH])
9         .....if (preprocess_vgg == 'Imagenet'):
10            .....image = preprocess_imagenet(image)
```

```

11 .....elif (preprocess_vgg == 'VGGFace'):
12 .....    # function does not accept tensor
13 .....    # so convert to np and then to tensor
14 .....    # image = np.array(image)
15 .....    image = preprocess_vggface(image, version=2)
16 .....    # image = tf.convert_to_tensor(image)
17 .....else:
18 .....    image = image / 255
19
20 .....return image
21
22     filenames = []
23     talk_regex = '*-01.jpg' if talk_frame else '*.jpg'
24     acted_regex = '02' if acted_frame else '*'
25     gen_regex = f'*-*-*{acted_regex}-*-*-*{talk_regex}'
26
27     file_dict = dict()
28     for label in sorted(label_names):
29         file_dict[label] = []
30
31
32     for label in label_names:
33         for actor in actors_idx:
34             for file in Path(path, label, actor).glob(f'{gen_regex}'):
35                 file_dict[label].append(str(file))
36
37
38     if shuffle:
39         for label, item in file_dict.items():
40             logging.info(f'Label: {label}')
41             logging.info(f'Array len: {len(item)}')
42             random.Random(SEED).shuffle(item)
43
44     arr_len = [len(arr) for arr in file_dict.values()]
45
46     if undersampling:
47         filenames = [arr[:min(arr_len)] for arr in file_dict.values()]
48     else:
49         filenames = [arr for arr in file_dict.values()]
50
51     filenames = sum(filenames, [])
52
53     if shuffle:
54         random.Random(SEED).shuffle(filenames)
55
56     labels = [
57         label_names.index(EMOTIONS_LABELS[int(utils.get_class(elem)) - 1])
58         for elem in filenames
59     ]
60
61     if (sampling < 1):
62         filenames, _, labels, _ = train_test_split(
63             filenames, labels, train_size=sampling, random_state=SEED
64         )
65
66     filenames_ds = tf.data.Dataset.from_tensor_slices(filenames)
67     labels_ds = tf.data.Dataset.from_tensor_slices(labels)
68
69     images_ds = filenames_ds.map(
70         parse_image, num_parallel_calls=tf.data.experimental.AUTOTUNE
71     )
72     ds = tf.data.Dataset.zip((images_ds, labels_ds))
73     # ds = configure_for_performance(ds)
74
75     return [ds, filenames]

```

```

1 sampling_rate = 1
2 talk_frame = True
3 acted_frame = False
4 preprocess_vgg = 'VGGFace' # False, Imagenet or VGGFace
5
6
7
8 train_ds, train_files = make_dataset(
9     const.frames_path, train_idx, talk_frame=talk_frame, acted_frame=acted_frame,
10     preprocess_vgg=preprocess_vgg, shuffle=True, sampling=sampling_rate
11 )
12
13 val_ds, val_files = make_dataset(
14     const.frames_path, val_idx, talk_frame=talk_frame, acted_frame=acted_frame,
15     preprocess_vgg=preprocess_vgg, sampling=sampling_rate
16 )
17
18 test_ds, test_files = make_dataset(
19     const.frames_path, test_idx, talk_frame=talk_frame, acted_frame=acted_frame,
20     preprocess_vgg=preprocess_vgg, sampling=sampling_rate
21 )

```

```

1 # train_ds = train_ds[:31882]
2 # train_files = train_files[:31882]
3
4 # val_ds = val_ds[:8426]
5 # val_files = val_files[:8426]
6
7 # test_ds = test_ds[:8301]
8 # test_files = test_files[:8301]

```

```

1 assert len(train_ds) == len(train_files), len(train_files)
2 assert len(val_ds) == len(val_files), len(val_files)
3 assert len(test_ds) == len(test_files), len(test_files)

```

```

1 train_ds_elements = len(train_ds)
2 test_ds_elements = len(test_ds)
3 val_ds_elements = len(val_ds)

```

```

1 print(f'train_ds samples: {train_ds_elements}')
2 print(f'test_ds samples: {test_ds_elements}')
3 print(f'val_ds samples: {val_ds_elements}')

```

```

🔄 train_ds samples: 67906
   test_ds samples: 17994
   val_ds samples: 17642

```

✓ Build and train the model

Add operations to reduce read latency while training the model:

`ds.batch` Combines consecutive elements of the dataset into batches. The components of the resulting element will have an additional outer dimension, which will be *batch_size*

`ds.cache` Caches the elements in this dataset.

`ds.prefetch` Allows later elements to be prepared while the current element is being processed. This often improves latency and throughput, at the cost of using additional memory to store prefetched elements.

```

1 def configure_for_performance(ds, batch_size=BATCH_SIZE):
2     ds = ds.batch(batch_size)
3     # ds = ds.cache()
4     # ds = ds.shuffle(buffer_size=1000)
5     # ds = ds.repeat()
6     ds = ds.prefetch(buffer_size=tf.data.AUTOTUNE)
7     return ds

```

```

1 train_ds = configure_for_performance(train_ds)
2 val_ds = configure_for_performance(val_ds)
3 test_ds = configure_for_performance(test_ds)

```

```

1 for example_images, example_labels in train_ds.take(1):
2     print(example_images.shape)
3     print(example_labels.shape)

```

```

🔄 (128, 224, 224, 3)
   (128,)

```

```


1 plt.figure(figsize=(16, 10))
2 rows = 2
3 cols = 2
4 n = rows * cols
5 for i in range(n):
6     plt.subplot(rows, cols, i + 1)
7     image = example_images[i]
8     plt.imshow(image) # 3 channels
9     # plt.imshow(image * 255, cmap='gray', vmin=0, vmax=255) # 1 channel
10    plt.title(f'{label_names[example_labels[i]]}\n{train_files[i]}')

```




neutral

Generated/Frames/neutral/11/01-01-01-01-02-01-11-035-01.jpg




sad

Generated/Frames/sad/15/01-01-04-02-01-01-15-068-01.jpg




disgust

Generated/Frames/disgust/04/01-01-07-01-01-01-04-071-01.jpg



fearful

Generated/Frames/fearful/13/01-01-06-02-02-13-031-01.jpg



```

1 def create_Bilotti_CNN(name='Bilotti_CNN'):
2
3     inputs = Input(shape=(IMG_HEIGHT, IMG_WIDTH, IMG_CHANNELS))
4
5     conv1 = Conv2D(32, kernel_size=(3, 3), activation='relu')(inputs)
6     conv2 = Conv2D(32, kernel_size=(3, 3), activation='relu')(conv1)
7     pool1 = MaxPooling2D(pool_size=(2, 2))(conv2)
8
9     conv3 = Conv2D(64, kernel_size=(3, 3), activation='relu')(pool1)
10    conv4 = Conv2D(64, kernel_size=(3, 3), activation='relu')(conv3)
11    pool2 = MaxPooling2D(pool_size=(2, 2))(conv4)
12
13    conv4 = Conv2D(18, kernel_size=(3, 3), activation='relu')(pool2)
14    conv5 = Conv2D(18, kernel_size=(3, 3), activation='relu')(conv4)
15    conv6 = Conv2D(18, kernel_size=(3, 3), activation='relu')(conv5)
16    pool3 = MaxPooling2D(pool_size=(2, 2))(conv6)
17
18    conv7 = Conv2D(56, kernel_size=(3, 3), activation='relu')(pool3)
19    conv8 = Conv2D(56, kernel_size=(3, 3), activation='relu')(conv7)
20    conv9 = Conv2D(56, kernel_size=(3, 3), activation='relu')(conv8)
21    pool4 = MaxPooling2D(pool_size=(2, 2))(conv9)
22
23    conv10 = Conv2D(51, kernel_size=(3, 3), activation='relu')(pool4)
24    conv11 = Conv2D(51, kernel_size=(3, 3), activation='relu')(conv10)
25    conv12 = Conv2D(51, kernel_size=(3, 3), activation='relu')(conv11)
26    pool5 = MaxPooling2D(pool_size=(2, 2))(conv12)
27
28    flatten = Flatten()(pool5)
29
30    dense1 = Dense(2048, activation='relu')(flatten)
31    drop1 = Dropout(0.25)(dense1)
32
33    dense2 = Dense(1024, activation='relu')(drop1)
34    drop2 = Dropout(0.4)(dense2)
35
36    output = Dense(NUM_CLASSES, activation='softmax')(drop2)
37
38    model = Model(inputs, output)
39
40    model._name = name
41
42    return model

```

```

1 def create_VGG16_Imagenet(name='VGG16_Imagenet'):
2
3     base_model = VGG16(
4         weights='imagenet',
5         include_top=False,
6         input_shape=(IMG_WIDTH, IMG_HEIGHT, IMG_CHANNELS)
7     )
8     base_model.trainable = False # Not trainable weights
9
10    flatten_layer = Flatten()
11    dense_layer_1 = Dense(2048, activation='relu')
12    drop_1 = Dropout(0.4)
13    dense_layer_2 = Dense(1024, activation='relu')
14    drop_2 = Dropout(0.4)
15    dense_layer_3 = Dense(512, activation='relu')
16    drop_3 = Dropout(0.4)
17    prediction_layer = Dense(NUM_CLASSES, activation='softmax')
18
19    model = Sequential([
20        base_model,
21        flatten_layer,
22        dense_layer_1,
23        drop_1,
24        dense_layer_2,
25        drop_2,
26        dense_layer_3,
27        drop_3,
28        prediction_layer
29    ])
30
31    model._name = name
32
33    return model

```

```

1 def create_VGG16_VGGFACE(name='VGG16_VGGFACE'):
2     nb_class = NUM_CLASSES
3
4     vgg_model = VGGFace(
5         include_top=False, weights='vggface', input_shape=(IMG_WIDTH, IMG_HEIGHT, IMG_CHANNELS)
6     )
7     last_layer = vgg_model.get_layer('pool5').output
8     x = Flatten(name='flatten')(last_layer)
9
10    x = Dense(512, activation='relu', name='fc6')(x)
11    x = Dropout(0.35)(x)
12    x = Dense(256, activation='relu', name='fc7')(x)
13    x = Dropout(0.35)(x)
14    x = Dense(128, activation='relu', name='fc8')(x)
15    x = Dropout(0.35)(x)
16
17    out = Dense(nb_class, activation='softmax', name='fc9')(x)
18
19    custom_vgg_model = Model(vgg_model.input, out)
20    custom_vgg_model._name = name
21
22    return custom_vgg_model

```

```

1 tf.keras.backend.clear_session() # clear all precedent models and sessions

```

```


1 check_path = 'checkpoint.weights.h5'
2 checkpointer = ModelCheckpoint(
3     check_path, monitor='val_accuracy', verbose=1, save_best_only=True,
4     save_weights_only=False, mode='auto', save_freq='epoch'
5 )

```

```

1 # model = create_cnn_model()
2 # model = medium_model()
3 # model = create_VGG16_Imagenet()
4 model = create_VGG16_VGGFACE()
5 # model = create_grigorasi_model()
6 # model = create_Bilotti_CNN()
7 model.summary()

```

 Downloading data from https://github.com/rcmalli/keras-vggface/releases/download/v2.0/rcmalli_vggface_tf_notop_vgg16.h5
58916864/58909280 [=====] - 0s 0us/step
Model: "VGG16_VGGFACE"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 224, 224, 3)	0
conv1_1 (Conv2D)	(None, 224, 224, 64)	1792
conv1_2 (Conv2D)	(None, 224, 224, 64)	36928
pool1 (MaxPooling2D)	(None, 112, 112, 64)	0
conv2_1 (Conv2D)	(None, 112, 112, 128)	73856
conv2_2 (Conv2D)	(None, 112, 112, 128)	147584
pool2 (MaxPooling2D)	(None, 56, 56, 128)	0
conv3_1 (Conv2D)	(None, 56, 56, 256)	295168
conv3_2 (Conv2D)	(None, 56, 56, 256)	590080
conv3_3 (Conv2D)	(None, 56, 56, 256)	590080
pool3 (MaxPooling2D)	(None, 28, 28, 256)	0
conv4_1 (Conv2D)	(None, 28, 28, 512)	1180160
conv4_2 (Conv2D)	(None, 28, 28, 512)	2359808
conv4_3 (Conv2D)	(None, 28, 28, 512)	2359808
pool4 (MaxPooling2D)	(None, 14, 14, 512)	0
conv5_1 (Conv2D)	(None, 14, 14, 512)	2359808
conv5_2 (Conv2D)	(None, 14, 14, 512)	2359808
conv5_3 (Conv2D)	(None, 14, 14, 512)	2359808
pool5 (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc6 (Dense)	(None, 512)	12845568
dropout (Dropout)	(None, 512)	0
fc7 (Dense)	(None, 256)	131328

dropout_1 (Dropout)	(None, 256)	0
fc8 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0

```
1 # Define training callbacks
2
3 class TimeHistory(tf.keras.callbacks.Callback):
4     def on_train_begin(self, logs={}):
5         self.times = []
6
7     def on_epoch_begin(self, batch, logs={}):
8         self.epoch_time_start = time.time()
9
10    def on_epoch_end(self, batch, logs={}):
11        self.times.append(time.time() - self.epoch_time_start)
12
13
14 early_stopping_callback = tf.keras.callbacks.EarlyStopping(
15     verbose=1,
16     patience=5,
17     restore_best_weights=True
18 )
19
20 reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, verbose=1,
21                               patience=3, min_lr=0)
```

```
1 METRICS = ['accuracy']
```

```
1 model.compile(
2     optimizer=tf.keras.optimizers.Adam(
3         learning_rate=1e-4
4     ),
5     # optimizer=tf.keras.optimizers.SGD(), # for VGG16_VGGFACE
6     loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
7     metrics=METRICS,
8 )
```

Train Model

```
1 EPOCHS = 100
2 time_callback = TimeHistory()
3 history = model.fit(
4     x=train_ds,
5     validation_data=val_ds,
6     epochs=EPOCHS,
7     callbacks=[time_callback, early_stopping_callback, reduce_lr]
8 )
```

Epoch 1/100
531/531 [=====] - 1139s 2s/step - loss: 0.9770 - accuracy: 0.6859 - val_loss: 2.4275 - val_accuracy: 0.4938 - lr: 1.0000e
Epoch 2/100
531/531 [=====] - 1009s 2s/step - loss: 0.1395 - accuracy: 0.9586 - val_loss: 2.2379 - val_accuracy: 0.5167 - lr: 1.0000e
Epoch 3/100
531/531 [=====] - 1008s 2s/step - loss: 0.0579 - accuracy: 0.9831 - val_loss: 2.6187 - val_accuracy: 0.5584 - lr: 1.0000e
Epoch 4/100
531/531 [=====] - 1004s 2s/step - loss: 0.0409 - accuracy: 0.9883 - val_loss: 2.2205 - val_accuracy: 0.5373 - lr: 1.0000e
Epoch 5/100
531/531 [=====] - 1002s 2s/step - loss: 0.0372 - accuracy: 0.9898 - val_loss: 2.8144 - val_accuracy: 0.5283 - lr: 1.0000e
Epoch 6/100
531/531 [=====] - 1004s 2s/step - loss: 0.0254 - accuracy: 0.9929 - val_loss: 3.1163 - val_accuracy: 0.5279 - lr: 1.0000e
Epoch 7/100
531/531 [=====] - ETA: 0s - loss: 0.0254 - accuracy: 0.9934
Epoch 7: ReduceLROnPlateau reducing learning rate to 1.9999999494757503e-05.
531/531 [=====] - 1004s 2s/step - loss: 0.0254 - accuracy: 0.9934 - val_loss: 2.9632 - val_accuracy: 0.4965 - lr: 1.0000e
Epoch 8/100
531/531 [=====] - 1001s 2s/step - loss: 0.0028 - accuracy: 0.9993 - val_loss: 3.2981 - val_accuracy: 0.5138 - lr: 2.0000e
Epoch 9/100
531/531 [=====] - ETA: 0s - loss: 5.5264e-04 - accuracy: 1.0000Restoring model weights from the end of the best epoch: 4.
531/531 [=====] - 1004s 2s/step - loss: 5.5264e-04 - accuracy: 1.0000 - val_loss: 3.7587 - val_accuracy: 0.5125 - lr: 2.0
Epoch 9: early stopping

```
1 EPOCHS = len(time_callback.times)
```

```

1 # Create model path
2 model_path = Path(const.models_path, 'Frame', model._name)
3 run_folders = list(Path(const.models_path, 'Frame', model._name).glob('Run_*'))
4
5 if not run_folders:
6     model_path = Path(model_path, 'Run_1')
7 else:
8     last_run = run_folders.pop()
9     last_run_idx = Path(last_run).name.split('_')[-1]
10    model_path = Path(model_path, f'Run_{int(last_run_idx) + 1}')
11
12 model_path.mkdir(parents=True, exist_ok=False)

```

```

1 # Save info on the indexes used for train, val and test
2 ds_info_path = Path(model_path, f'{model._name}_dataset.txt')
3 with open(ds_info_path, 'w+', newline='') as res_file:
4     res_file.write(f'Train indexes: {train_idx}\n')
5     res_file.write(f'Train files: {train_ds_elements}\n')
6     res_file.write(f'Val indexes: {val_idx}\n')
7     res_file.write(f'Val files: {val_ds_elements}\n')
8     res_file.write(f'Test indexes: {test_idx}\n')
9     res_file.write(f'Test files: {test_ds_elements}\n')

```

```

1 metrics = history.history

```

```


1 mod_loss = metrics['loss']
2 mod_val_loss = metrics['val_loss']
3 mod_accuracy = metrics['accuracy']
4 mod_val_accuracy = metrics['val_accuracy']
5 # mod_f1 = metrics['fBeta_score']
6 # mod_val_f1 = metrics['val_fBeta_score']
7
8 mod_mean_loss = np.mean(mod_loss)
9 mod_mean_val_loss = np.mean(mod_val_loss)
10 mod_mean_accuracy = np.mean(mod_accuracy)
11 mod_mean_val_accuracy = np.mean(mod_val_accuracy)
12 # mod_mean_f1 = np.mean(mod_f1)
13 # mod_mean_val_f1 = np.mean(mod_val_f1)

```

```

1 # Save Loss
2 plt.title('Model loss')
3 plt.plot(history.epoch, mod_loss, mod_val_loss)
4 plt.legend(['loss', 'val_loss'])
5 plt.ylim([0, max(plt.ylim())])
6 plt.xlabel('Epoch')
7 plt.ylabel('Loss [CrossEntropy]')
8 plt.savefig(Path(model_path, 'loss.png'))


```



```

1 # Save Accuracy
2 plt.title('Model accuracy')
3 plt.plot(
4     history.epoch,
5     100 * np.array(mod_accuracy),
6     100 * np.array(mod_val_accuracy)
7 )
8 plt.legend(['accuracy', 'val_accuracy'])
9 plt.ylim([0, 100])
10 plt.xlabel('Epoch')
11 plt.ylabel('Accuracy [%]')
12 plt.savefig(Path(model_path, 'accuracy.png'))

```



Unsupported Cell Type. Double-Click to inspect/edit the content.

Unsupported Cell Type. Double-Click to inspect/edit the content.

Evaluate Model

```

1 model_eval = model.evaluate(test_ds, return_dict=True)

```

141/141 [=====] - 103s 730ms/step - loss: 2.1810 - accuracy: 0.5296

```

1 # Save model
2 model.save(Path(model_path, f'{model._name}.keras'), overwrite=False)
3 # Save history
4 np.save(Path(model_path, f'{model._name}_history.npy'), history)
5 # Save model image
6 model_img = tf.keras.utils.plot_model(
7     model, Path(model_path, f'{model._name}.png'), show_shapes=True,
8     show_layer_names=True, show_layer_activations=True
9 )

```

```

1 model_eval

```

{'loss': 2.181023120880127, 'accuracy': 0.5296210050582886}

```

1 test_loss = model_eval['loss']
2 test_accuracy = model_eval['accuracy']
3 # test_f1 = model_eval['fBeta_score']
4 mean_epoch_time = np.mean(time_callback.times)

```

```

1 # Salvataggio informazioni modello
2 model_save_path = Path(model_path, f'{model._name}_result.txt')
3 with open(model_save_path, 'w+', newline='') as res_file:
4     res_file.write(f'BATCH: {BATCH_SIZE}\n')
5     res_file.write(f'Train loss: {str(mod_loss)}\n')
6     res_file.write(f'val_loss: {str(mod_val_loss)}\n')
7     res_file.write(f'Train accuracy: {str(mod_accuracy)}\n')
8     res_file.write(f'Train val_accuracy: {str(mod_val_accuracy)}\n')
9     # res_file.write(f'Train f1_score: {str(mod_f1)}\n')
10    # res_file.write(f'Train val_f1_score: {str(mod_val_f1)}\n')
11    res_file.write(f'Test loss: {str(test_loss)}\n')
12    res_file.write(f'Test accuracy: {str(test_accuracy)}\n')
13    # res_file.write(f'Test f1_score: {str(test_f1)}\n')
14    res_file.write(f'Mean epoch time: {str(mean_epoch_time)}')

```

```

1 # Salvataggio informazioni generali modelli
2 with open(Path(const.models_path, 'Frame', 'models.csv'), 'a+') as csvfile:
3     filewriter = csv.writer(
4         csvfile, delimiter=';', quotechar='|', quoting=csv.QUOTE_MINIMAL
5     )
6
7     # filewriter.writerow(
8     #     ["Model Name", "Epochs", "% Validation", "% Test set",
9     #     "Train loss", "Train accuracy", "Val loss", "Val accuracy",
10    #     "Test loss", "Test accuracy", "Mean epoch time", "Note"]
11    # )
12    test_ds_perc = utils.trunc((test_ds_elements * 100) / TOTAL_ELEMENTS, 2)
13    val_ds_perc = utils.trunc((val_ds_elements * 100) / TOTAL_ELEMENTS, 2)
14    full_path = str(Path(model._name, model_path.name))
15    filewriter.writerow(
16        [full_path, EPOCHS, val_ds_perc, test_ds_perc,
17         mod_loss, mod_accuracy, mod_val_loss, mod_val_accuracy,
18         test_loss, test_accuracy, mean_epoch_time, '']
19    )

```

```

1 for test_images, test_labels in test_ds.take(1):
2     print(test_images.shape)
3     print(test_labels.shape)

```

```

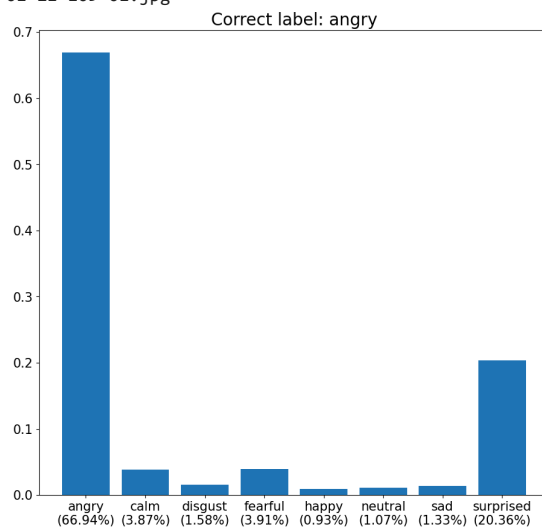

➡ (128, 224, 224, 3)
(128,)

```

```

1 gen = np.random.default_rng(seed=None)
2 idx = gen.integers(0, len(test_images))
3 print(test_files[idx])
4
5 image = test_images[idx]
6 label = test_labels[idx]
7
8 net_input = utils.extend_tensor(image, 0)
9 prediction = model(net_input)
10 prediction = prediction[0].numpy()
11
12 valued_arr = []
13
14 for idx, name in enumerate(label_names):
15     valued_arr.append(f'{name}\n({prediction[idx]:.2%})')
16
17 fig, ax = plt.subplots(
18     nrows=1, ncols=2, width_ratios=[0.4, 0.6], figsize=(20, 10)
19 )
20
21 pltDisplay(image * 255, ax=ax[0]) # 1 channel
22 # pltDisplay(image, ax=ax[0])
23
24 ax[1].bar(valued_arr, prediction)
25 plt.xticks(fontsize=15)
26 plt.yticks(fontsize=15)
27 plt.title(f'Correct label: {label_names[label]}', fontsize=20)
28 # plt.xlabel('Predicted class')
29 # plt.ylabel('Percentage')
30 plt.show()

```



Display a confusion matrix

Use a [confusion matrix](#) to check how well the model did classifying each of the commands in the test set:

Unsupported Cell Type. Double-Click to inspect/edit the content.

```
1 y_pred = model.predict(test_ds)
```

141/141 [=====] - 96s 680ms/step

```
1 y_pred = tf.argmax(y_pred, axis=1, output_type=tf.int32)
```

```
1 y_true = tf.concat(list(test_ds.map(lambda _, lab: lab)), axis=0)
```


```
1 report = classification_report(  
2     y_true, y_pred, target_names=label_names,  
3     output_dict=True, zero_division='warn'  
4 )
```

```
1 rep_to_csv = pd.DataFrame(data=report).transpose()
```

```

1 fig, ax = plt.subplots(figsize=(10, 10))
2 ax.xaxis.tick_top()
3 sns.heatmap(rep_to_csv.iloc[:NUM_CLASSES, :3],
4             cbar=True,
5             square=False,
6             annot=True,
7             annot_kws={'size': 15},
8             fmt='.2g',
9             linewidths=0.5)
10 plt.savefig(Path(model_path, f'{model._name}_f1_score.png'))
11 plt.show()
12
13
14 with sns.axes_style('white'):
15     fig, ax = plt.subplots(figsize=(3, 5))
16     ax.xaxis.tick_top()
17     sns.heatmap(rep_to_csv.iloc[:NUM_CLASSES, 3:],
18               cbar=False,
19               square=False,
20               annot=True,
21               annot_kws={'size': 15},
22               fmt='.4g',
23               cmap=ListedColormap([('purple'])),
24               linewidths=0.5)
25     plt.savefig(Path(model_path, f'{model._name}_support.png'))
26     plt.show()

```



```
1 report_save_path = Path(model_path, f'{model._name}_report.csv')
2 rep_to_csv.to_csv(report_save_path)
```

Unsupported Cell Type. Double-Click to inspect/edit the content.


```
1 np.mean([report[c]['f1-score'] for c in list(report)[:NUM_CLASSES]])
```

0.523185237626927

```

1 # True Test Distribution
2 # unique, counts = np.unique(y_true, return_counts=True)
3 # collections.Counter(y_true)
4 counts = [np.count_nonzero(y_true == idx) for idx in range(len(label_names))]
5 valued_arr = []
6 # for i in range(len(label_names)):
7 for idx, name in enumerate(label_names):
8     count = counts[idx]
9     valued_arr.append(f'{name}\n({count})')
10
11 fig = plt.subplots(figsize=(20, 10))
12 plt.bar(valued_arr, counts)
13 plt.xticks(fontsize=15)
14 plt.yticks(fontsize=15)
15 plt.title('Test True Distribution', fontsize=20)
16 # plt.xlabel('Predicted class')
17 # plt.ylabel('Percentage')
18 plt.savefig(Path(model_path, f'{model._name}_trueDist.png'))
19 plt.show()

```




```

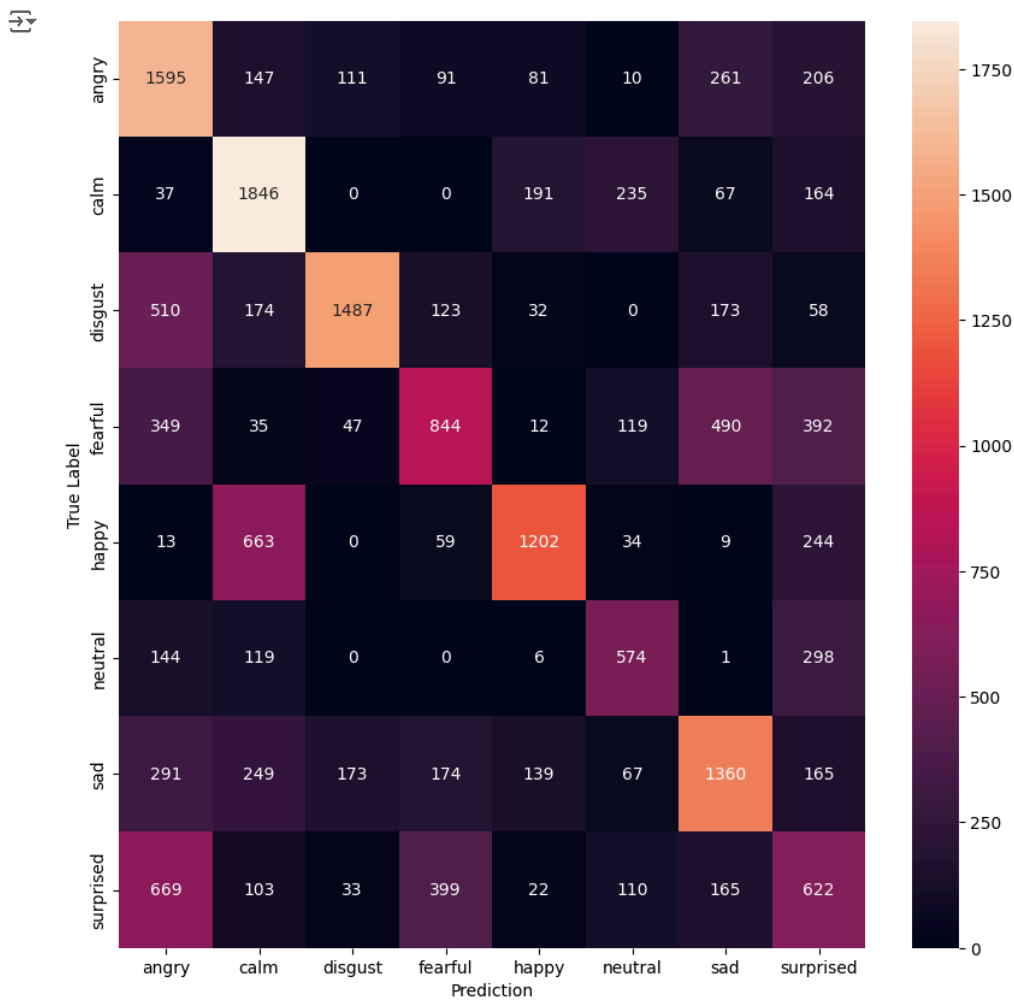
1 # Predicted Test Distribution
2 # unique, counts = np.unique(y_pred, return_counts=True)
3 counts = [np.count_nonzero(y_pred == idx) for idx in range(len(label_names))]
4 valued_arr = []
5 unique_idx = 0
6 for idx, name in enumerate(label_names):
7     count = counts[idx]
8     valued_arr.append(f'{name}\n({count})')
9
10 fig = plt.subplots(figsize=(20, 10))
11 plt.bar(valued_arr, counts)
12 plt.xticks(fontsize=15)
13 plt.yticks(fontsize=15)
14 plt.title('Test result distribution', fontsize=20)
15 # plt.xlabel('Predicted class')
16 # plt.ylabel('Percentage')
17 plt.savefig(Path(model_path, f'{model._name}_predDist.png'))
18 plt.show()

```


Test result distribution



```
1 confusion_mtx = tf.math.confusion_matrix(y_true, y_pred)
2 fig, ax = plt.subplots(figsize=(10, 10))
3 sns.heatmap(confusion_mtx,
4             xticklabels=label_names,
5             yticklabels=label_names,
6             annot=True, fmt='.4g')
7 plt.xlabel('Prediction')
8 plt.ylabel('True Label')
9 plt.savefig(Path(model_path, f'{model._name}_heat.png'))
10 plt.show()
```



```

1 def multiclass_roc_auc_score(target, y_test, y_pred, average="macro"):
2     # function for scoring roc auc score for multi-class
3     lb = LabelBinarizer()
4     lb.fit(y_test)
5     y_test = lb.transform(y_test)
6     y_pred = lb.transform(y_pred)
7
8     if len(target) > 2:
9         for (idx, c_label) in enumerate(target):
10
11             fpr, tpr, thresholds = roc_curve(
12                 y_test[:, idx].astype(int),
13                 y_pred[:, idx]
14             )
15             c_ax.plot(
16                 fpr, tpr, label='%s (AUC:%0.2f)' % (c_label, auc(fpr, tpr))
17             )
18     else:
19         fpr, tpr, thresholds = roc_curve(
20             y_test,
21             y_pred
22         )
23         c_ax.plot(
24             fpr, tpr, label='Model (AUC:%0.2f)' % (auc(fpr, tpr)), color='#ff7f0e'
25         )
26     c_ax.plot(fpr, fpr, color='b', linestyle='--', label='Random Guessing')
27
28     return roc_auc_score(y_test, y_pred, average=average)


```

```

1 # set plot figure size
2 fig, c_ax = plt.subplots(1, 1, figsize=(12, 8))
3
4 print('ROC AUC score:', multiclass_roc_auc_score(
5     label_names,
6     tf.reshape(y_true, (y_true.shape[0], 1)),
7     tf.reshape(y_pred, (y_pred.shape[0], 1))
8 ))
9
10 c_ax.legend()
11 c_ax.set_xlabel('False Positive Rate')
12 c_ax.set_ylabel('True Positive Rate')
13 plt.savefig(Path(model_path, f'{model._name}_ROC.png'))
14 plt.show()

```

ROC AUC score: 0.7268530424410194



Run inference on an image file

Finally, verify the model's prediction output using an image

```

1 def parse_image(filename):
2     image = tf.io.read_file(filename)
3     image = tf.image.decode_jpeg(image, channels=IMG_CHANNELS)
4     image = tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH])
5     if (preprocess_vgg == 'Imagenet'):
6         image = preprocess_imagenet(image)
7     elif (preprocess_vgg == 'VGGFace'):
8         # function does not accept tensor
9         # so convert to np and then to tensor
10        # image = np.array(image)
11        image = preprocess_vggface(image, version=2)
12        # image = tf.convert_to_tensor(image)
13    else:
14        image = image / 255
15

```