

Import frames and train, test, evaluate network on facial images

✧ Importing Dependencies

```
1 from google.colab import drive
2 drive.mount('/content/drive', force_remount=True)
```

↳ Mounted at /content/drive

```
1 PROJECT_FOLDER = '/content/drive/MyDrive/CNN_Emotion_Classification/'
```

```
1 # %cd /content/drive/MyDrive/CNN_Emotion_Classification/
```

```
1 # copy content in main folder
2 ! cp -a {PROJECT_FOLDER}. ./
```

```
1 !pip install -r Requirements/pireqs_opencv_contrib_env.txt
```

↳ Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from -r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: opencv_contrib_python in /usr/local/lib/python3.10/dist-packages (from -r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from -r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: scikit_learn in /usr/local/lib/python3.10/dist-packages (from -r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from -r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (from -r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: tensorflow in /usr/local/lib/python3.10/dist-packages (from -r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: setuptools>=18.5 in /usr/local/lib/python3.10/dist-packages (from ipython->-r Requirements/pireqs_opencv_contrib_env.txt)

Collecting jedi>=0.16 (from ipython->-r Requirements/pireqs_opencv_contrib_env.txt (line 2))

Downloading jedi-0.19.1-py2.py3-none-any.whl (1.6 MB)

1.6/1.6 MB 31.1 MB/s eta 0:00:00

Requirement already satisfied: decorator in /usr/local/lib/python3.10/dist-packages (from ipython->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: pickleshare in /usr/local/lib/python3.10/dist-packages (from ipython->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: traitlets>=4.2 in /usr/local/lib/python3.10/dist-packages (from ipython->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: prompt-toolkit!=3.0.0,!<3.0.1,<3.1.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from ipython->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: pygments in /usr/local/lib/python3.10/dist-packages (from ipython->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: backcall in /usr/local/lib/python3.10/dist-packages (from ipython->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: matplotlib-inline in /usr/local/lib/python3.10/dist-packages (from ipython->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: pexpect>4.3 in /usr/local/lib/python3.10/dist-packages (from ipython->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: h5py in /usr/local/lib/python3.10/dist-packages (from keras_vggface->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: pillow in /usr/local/lib/python3.10/dist-packages (from keras_vggface->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: keras in /usr/local/lib/python3.10/dist-packages (from keras_vggface->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: six>=1.9.0 in /usr/local/lib/python3.10/dist-packages (from keras_vggface->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: pyyaml in /usr/local/lib/python3.10/dist-packages (from keras_vggface->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: audioread>=2.1.9 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: joblib>=0.14 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: numba>=0.51.0 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: soundfile>=0.12.1 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: pooch>=1.1 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: soxr>=0.3.2 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: typing-extensions>=4.1.1 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: lazy-loader>=0.1 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: msgpack>=1.0 in /usr/local/lib/python3.10/dist-packages (from librosa->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit_learn->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: flatbuffers>=23.5.26 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: gast!=0.5.0,!<0.5.1,!<0.5.2,>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: ml-dtypes<=0.2.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: protobuf!=4.21.0,!<4.21.1,!<4.21.2,!<4.21.3,!<4.21.4,!<4.21.5,<5.0.0dev,>=3.20.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: wrapt<1.15,>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: tensorflowboard<2.16,>=2.15.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Requirement already satisfied: tensorflow-estimator<2.16,>=2.15.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow->-r Requirements/pireqs_opencv_contrib_env.txt)

Double-click (or enter) to edit

```
1
```

```
1 # !pip install keras_applications
```

```
1 from display import pltDisplay
2 from pathlib import Path
3 from matplotlib.colors import ListedColormap
4 from sklearn.metrics import classification_report, roc_curve, auc, roc_auc_score
5 from sklearn.preprocessing import LabelBinarizer
6 from sklearn.model_selection import train_test_split
7 # from sorting import human_sort
8 from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, Dense, Flatten
9 from tensorflow.keras.layers import BatchNormalization, Dropout, Activation, ReLU, Softmax
10 from tensorflow.keras.models import Model, Sequential
11 from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau
12 from tensorflow.keras.regularizers import L2
13 # import IPython.display as ipd
14
15
16 import constants as const
17 import csv
18 import cv2
19 import json
20 import logging
21 import gc
22 import matplotlib.pyplot as plt
23 import numpy as np
24 import pandas as pd
25 import random
26 import seaborn as sns
27 import tensorflow as tf
28 import subprocess
29 import tensorflow.keras.backend as K
30 import time
31 import utils
32
33 from keras_vggface.vggface import VGGFace
34 from tensorflow.keras.applications.vgg16 import VGG16
35 from tensorflow.keras.applications.vgg16 import preprocess_input as preprocess_imagenet
36 from keras_vggface.utils import preprocess_input as preprocess_vggface
```

```
1 print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
```

```
➦ Num GPUs Available: 1
```

```
1 gpus = tf.config.list_physical_devices('GPU')
2 if gpus:
3     try:
4         # Currently, memory growth needs to be the same across GPUs
5         for gpu in gpus:
6             print("Name:", gpu.name, " Type:", gpu.device_type)
7             print(tf.config.experimental.get_device_details(gpu))
8             print(tf.config.experimental.get_memory_info('GPU:0'))
9             # tf.config.experimental.set_memory_growth(gpu, True)
10        logical_gpus = tf.config.list_logical_devices('GPU')
11        print(len(gpus), "Physical GPUs,", len(logical_gpus), "Logical GPUs")
12    except RuntimeError as e:
13        # Memory growth must be set before GPUs have been initialized
14        print(e)
```

```
➦ Name: /physical_device:GPU:0 Type: GPU
{'compute_capability': (7, 5), 'device_name': 'Tesla T4'}
{'current': 0, 'peak': 0}
1 Physical GPUs, 1 Logical GPUs
```

```
1 # unzip archive in folder
2 !7z x Generated/Frames_300.zip -oGenerated/
```

```
➦
7-Zip [64] 16.02 : Copyright (c) 1999-2016 Igor Pavlov : 2016-05-21
p7zip Version 16.02 (locale=en-US.UTF-8,Utf16=on,HugeFiles=on,64 bits,2 CPUs Intel(R) Xeon(R) CPU @ 2.00GHz (50653),ASM,AES-NI)

Scanning the drive for archives:
1 file, 4154172627 bytes (3962 MiB)

Extracting archive: Generated/Frames_300.zip
--
Path = Generated/Frames_300.zip
Type = zip
Physical Size = 4154172627
```

```
64-bit = +  
  
Everything is Ok  
  
Folders: 202  
Files: 158556  
Size: 4143561934  
Compressed: 4154172627
```

```
1 utils.modules_info()
```



```
OpenCV:  
  Version: 4.8.0  
Tensorflow:  
  Version: 2.15.0
```

```
1 log_file = Path(const.logs_path, 'FACER.log')  
2 logging.basicConfig(  
3     format='%(asctime)s %(message)s',  
4     filemode='a',  
5     filename=log_file,  
6     encoding='utf-8',  
7     level=logging.INFO,  
8     force=True  
9 )
```

✓ Importing Dataset

```
1 data_df = pd.read_csv(Path(const.csv_path, 'dataset.csv'))
```

✓ Frame

✓ Preparing Data

✓ Dataset Creation for ML

```
1 IMG_WIDTH = IMG_HEIGHT = 224  
2 IMG_CHANNELS = 3  
3 SEED = 42  
4 BATCH_SIZE = 128  
5 VALIDATION_SPLIT = 0.2  
6 EMOTIONS_LABELS = const.EMOTIONS_LABELS # RAVDESS emotion labels
```

```
1 EMOTIONS_LABELS
```



```
['neutral', 'calm', 'happy', 'sad', 'angry', 'fearful', 'disgust', 'surprised']
```

```
1 TOTAL_ELEMENTS = const.DATASET_TOTAL_ELEMENTS  
2 label_names = const.EMOTIONS_LABELS_SORTED.copy()  
3 # label_names.remove('neutral')  
4 # label_names.remove('calm')  
5 # label_names.remove('surprised')  
6  
7 NUM_CLASSES = len(label_names)
```

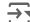
✓ Dataset Creation - NEW

```
1 actors_labels = [f'{i:02d}' for i in range(1, 25)]  
2 dist_idxes = {  
3     '1': [slice(0, 16), slice(16, 20), slice(20, 24)],  
4     '2': [slice(8, 24), slice(4, 8), slice(0, 4)],  
5     '3': [slice(4, 20), slice(20, 24), slice(0, 4)]  
6 }
```

```

1 # Split actors in train, validation, test
2 dist_n = 1
3 train_idx, val_idx, test_idx = [actors_labels[i] for i in dist_idx[str(dist_n)]]
4
5 print(train_idx, val_idx, test_idx)

```

 ['01', '02', '03', '04', '05', '06', '07', '08', '09', '10', '11', '12', '13', '14', '15', '16'] ['17', '18', '19', '20'] ['21', '22', '23', '24', '25', '26', '27', '28', '29', '30', '31', '32', '33', '34', '35', '36', '37', '38', '39', '40', '41', '42', '43', '44', '45', '46', '47', '48', '49', '50', '51', '52', '53', '54', '55', '56', '57', '58', '59', '60', '61', '62', '63', '64', '65', '66', '67', '68', '69', '70', '71', '72', '73', '74', '75', '76', '77', '78', '79', '80', '81', '82', '83', '84', '85', '86', '87', '88', '89', '90', '91', '92', '93', '94', '95', '96', '97', '98', '99', '100']

```

1 def make_dataset(path, actors_idx, talk_frame=False, acted_frame=False,
2                 undersampling=False, preprocess_vgg=True, shuffle=False, sampling=1):
3
4     def parse_image(filename):
5         image = tf.io.read_file(filename)
6         image = tf.image.decode_jpeg(image, channels=IMG_CHANNELS)
7         image = tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH])
8         if (preprocess_vgg):
9             image = preprocess_vgg(image)
10        else:
11            image = image / 255
12
13        return image
14
15    filenames = []
16    talk_regex = '*-01.jpg' if talk_frame else '*.jpg'
17    acted_regex = '*02' if acted_frame else '*'
18    gen_regex = f'*-*-*-{acted_regex}-*-*-*-{talk_regex}'
19
20    file_dict = dict()
21    for label in sorted(label_names):
22        file_dict[label] = []
23
24
25    for label in label_names:
26        for actor in actors_idx:
27            for file in Path(path, label, actor).glob(f'{gen_regex}'):
28                file_dict[label].append(str(file))
29
30
31    if shuffle:
32        for label, item in file_dict.items():
33            logging.info(f'Label: {label}')
34            logging.info(f'Array len: {len(item)}')
35            random.Random(SEED).shuffle(item)
36
37    arr_len = [len(arr) for arr in file_dict.values()]
38
39    if undersampling:
40        filenames = [arr[:min(arr_len)] for arr in file_dict.values()]
41    else:
42        filenames = [arr for arr in file_dict.values()]
43
44    filenames = sum(filenames, [])
45
46    if shuffle:
47        random.Random(SEED).shuffle(filenames)
48
49    labels = [
50        label_names.index(EMOTIONS_LABELS[int(utils.get_class(elem)) - 1])
51        for elem in filenames
52    ]
53
54    if (sampling < 1):
55        filenames, _, labels, _ = train_test_split(
56            filenames, labels, train_size=sampling, random_state=SEED
57        )
58
59    filenames_ds = tf.data.Dataset.from_tensor_slices(filenames)
60    labels_ds = tf.data.Dataset.from_tensor_slices(labels)
61
62    images_ds = filenames_ds.map(
63        parse_image, num_parallel_calls=tf.data.experimental.AUTOTUNE
64    )
65    ds = tf.data.Dataset.zip((images_ds, labels_ds))
66    # ds = configure_for_performance(ds)
67
68    return [ds, filenames]

```

```

1 sampling_rate = 1
2 talk_frame = True
3 acted_frame = False
4 preprocess_vgg = 'Imagenet' # False, Imagenet or VGGFace
5
6 if(preprocess_vgg == 'Imagenet'):
7     preprocess_vgg = preprocess_imagenet
8 elif(preprocess_vgg == 'VGGFace'):
9     preprocess_vgg = preprocess_vggface
10
11 train_ds, train_files = make_dataset(
12     const.frames_path, train_idx, talk_frame=talk_frame, acted_frame=acted_frame,
13     preprocess_vgg=preprocess_vgg, shuffle=True, sampling=sampling_rate
14 )
15
16 val_ds, val_files = make_dataset(
17     const.frames_path, val_idx, talk_frame=talk_frame, acted_frame=acted_frame,
18     preprocess_vgg=preprocess_vgg, sampling=sampling_rate
19 )
20
21 test_ds, test_files = make_dataset(
22     const.frames_path, test_idx, talk_frame=talk_frame, acted_frame=acted_frame,
23     preprocess_vgg=preprocess_vgg, sampling=sampling_rate
24 )

```

```

1 # train_ds = train_ds[:31882]
2 # train_files = train_files[:31882]
3
4 # val_ds = val_ds[:8426]
5 # val_files = val_files[:8426]
6
7 # test_ds = test_ds[:8301]
8 # test_files = test_files[:8301]

```

```

1 assert len(train_ds) == len(train_files), len(train_files)
2 assert len(val_ds) == len(val_files), len(val_files)
3 assert len(test_ds) == len(test_files), len(test_files)

```

```

1 train_ds_elements = len(train_ds)
2 test_ds_elements = len(test_ds)
3 val_ds_elements = len(val_ds)

```

```

1 print(f'train_ds samples: {train_ds_elements}')
2 print(f'test_ds samples: {test_ds_elements}')
3 print(f'val_ds samples: {val_ds_elements}')

```

```

↗ train_ds samples: 67906
  test_ds samples: 17994
  val_ds samples: 17642

```

✓ Build and train the model

Add operations to reduce read latency while training the model:

`ds.batch` Combines consecutive elements of the dataset into batches. The components of the resulting element will have an additional outer dimension, which will be *batch_size*

`ds.cache` Caches the elements in this dataset.

`ds.prefetch` Allows later elements to be prepared while the current element is being processed. This often improves latency and throughput, at the cost of using additional memory to store prefetched elements.

```

1 def configure_for_performance(ds, batch_size=BATCH_SIZE):
2     ds = ds.batch(batch_size)
3     # ds = ds.cache()
4     # ds = ds.shuffle(buffer_size=1000)
5     # ds = ds.repeat()
6     ds = ds.prefetch(buffer_size=tf.data.AUTOTUNE)
7     return ds

```

```

1 train_ds = configure_for_performance(train_ds)
2 val_ds = configure_for_performance(val_ds)
3 test_ds = configure_for_performance(test_ds)

```

```

1 for example_images, example_labels in train_ds.take(1):
2     print(example_images.shape)
3     print(example_labels.shape)

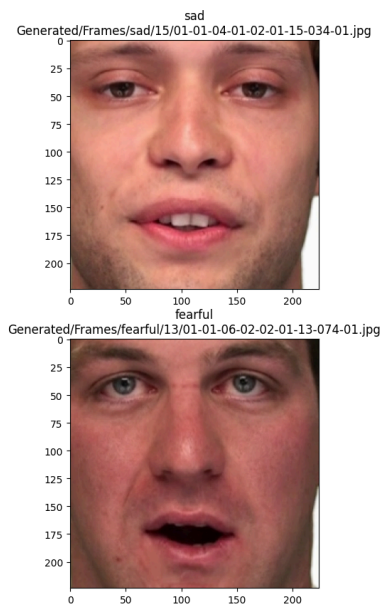
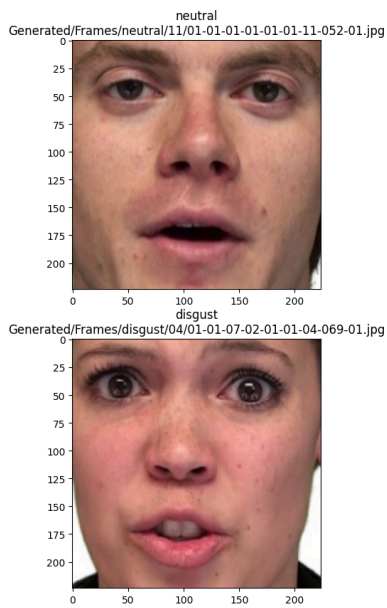
```

→ (128, 224, 224, 3)
(128,)

```

1 plt.figure(figsize=(16, 10))
2 rows = 2
3 cols = 2
4 n = rows * cols
5 for i in range(n):
6     plt.subplot(rows, cols, i + 1)
7     image = example_images[i]
8     plt.imshow(image) # 3 channels
9     # plt.imshow(image * 255, cmap='gray', vmin=0, vmax=255) # 1 channel
10    plt.title(f'{label_names[example_labels[i]]}\n{train_files[i]}')

```



```

1 def create_Bilotti_CNN(name='Bilotti_CNN'):
2
3     inputs = Input(shape=(IMG_HEIGHT, IMG_WIDTH, IMG_CHANNELS))
4
5     conv1 = Conv2D(32, kernel_size=(3, 3), activation='relu')(inputs)
6     conv2 = Conv2D(32, kernel_size=(3, 3), activation='relu')(conv1)
7     pool1 = MaxPooling2D(pool_size=(2, 2))(conv2)
8
9     conv3 = Conv2D(64, kernel_size=(3, 3), activation='relu')(pool1)
10    conv4 = Conv2D(64, kernel_size=(3, 3), activation='relu')(conv3)
11    pool2 = MaxPooling2D(pool_size=(2, 2))(conv4)
12
13    conv4 = Conv2D(18, kernel_size=(3, 3), activation='relu')(pool2)
14    conv5 = Conv2D(18, kernel_size=(3, 3), activation='relu')(conv4)
15    conv6 = Conv2D(18, kernel_size=(3, 3), activation='relu')(conv5)
16    pool3 = MaxPooling2D(pool_size=(2, 2))(conv6)
17
18    conv7 = Conv2D(56, kernel_size=(3, 3), activation='relu')(pool3)
19    conv8 = Conv2D(56, kernel_size=(3, 3), activation='relu')(conv7)
20    conv9 = Conv2D(56, kernel_size=(3, 3), activation='relu')(conv8)
21    pool4 = MaxPooling2D(pool_size=(2, 2))(conv9)
22
23    conv10 = Conv2D(51, kernel_size=(3, 3), activation='relu')(pool4)
24    conv11 = Conv2D(51, kernel_size=(3, 3), activation='relu')(conv10)
25    conv12 = Conv2D(51, kernel_size=(3, 3), activation='relu')(conv11)
26    pool5 = MaxPooling2D(pool_size=(2, 2))(conv12)
27
28    flatten = Flatten()(pool5)
29
30    dense1 = Dense(2048, activation='relu')(flatten)
31    drop1 = Dropout(0.25)(dense1)
32
33    dense2 = Dense(1024, activation='relu')(drop1)
34    drop2 = Dropout(0.4)(dense2)
35
36    output = Dense(NUM_CLASSES, activation='softmax')(drop2)
37
38    model = Model(inputs, output)
39
40    model._name = name
41
42    return model

```

```

1 def create_VGG16_Imagenet(name='VGG16_Imagenet'):
2
3     base_model = VGG16(
4         weights='imagenet',
5         include_top=False,
6         input_shape=(IMG_WIDTH, IMG_HEIGHT, IMG_CHANNELS)
7     )
8     base_model.trainable = False # Not trainable weights
9
10    flatten_layer = Flatten()
11    dense_layer_1 = Dense(2048, activation='relu')
12    drop_1 = Dropout(0.4)
13    dense_layer_2 = Dense(1024, activation='relu')
14    drop_2 = Dropout(0.4)
15    dense_layer_3 = Dense(512, activation='relu')
16    drop_3 = Dropout(0.4)
17    prediction_layer = Dense(NUM_CLASSES, activation='softmax')
18
19    model = Sequential([
20        base_model,
21        flatten_layer,
22        dense_layer_1,
23        drop_1,
24        dense_layer_2,
25        drop_2,
26        dense_layer_3,
27        drop_3,
28        prediction_layer
29    ])
30
31    model._name = name
32
33    return model

```

```

1 def create_VGG16_VGGFACE(name='VGG16_VGGFACE'):
2     nb_class = NUM_CLASSES
3
4     vgg_model = VGGFace(
5         include_top=False, weights='vggface', input_shape=(IMG_WIDTH, IMG_HEIGHT, IMG_CHANNELS)
6     )
7     last_layer = vgg_model.get_layer('pool5').output
8     x = Flatten(name='flatten')(last_layer)
9
10    x = Dense(512, activation='relu', name='fc6')(x)
11    x = Dropout(0.35)(x)
12    x = Dense(256, activation='relu', name='fc7')(x)
13    x = Dropout(0.35)(x)
14    x = Dense(128, activation='relu', name='fc8')(x)
15    x = Dropout(0.35)(x)
16
17    out = Dense(nb_class, activation='softmax', name='fc9')(x)
18
19    custom_vgg_model = Model(vgg_model.input, out)
20    custom_vgg_model._name = name
21
22    return custom_vgg_model

```

```

1 tf.keras.backend.clear_session() # clear all precedent models and sessions

```

```

1 check_path = 'checkpoint.weights.h5'
2 checkpointer = ModelCheckpoint(
3     check_path, monitor='val_accuracy', verbose=1, save_best_only=True,
4     save_weights_only=False, mode='auto', save_freq='epoch'
5 )

```

```

1 # model = create_cnn_model()
2 # model = medium_model()
3 model = create_VGG16_Imagenet()
4 # model = create_grigorasi_model()
5 # model = create_Bilotti_CNN()
6 model.summary()

```

➞ Model: "VGG16_Imagenet"

Layer (type)	Output Shape	Param #
=====		
vgg16 (Functional)	(None, 7, 7, 512)	14714688
flatten_2 (Flatten)	(None, 25088)	0
dense_7 (Dense)	(None, 2048)	51382272
dropout_5 (Dropout)	(None, 2048)	0
dense_8 (Dense)	(None, 1024)	2098176
dropout_6 (Dropout)	(None, 1024)	0
dense_9 (Dense)	(None, 512)	524800
dropout_7 (Dropout)	(None, 512)	0
dense_10 (Dense)	(None, 8)	4104
=====		
Total params: 68724040 (262.16 MB)		
Trainable params: 54009352 (206.03 MB)		
Non-trainable params: 14714688 (56.13 MB)		


```

1 # Define training callbacks
2
3 class TimeHistory(tf.keras.callbacks.Callback):
4     def on_train_begin(self, logs={}):
5         self.times = []
6
7     def on_epoch_begin(self, batch, logs={}):
8         self.epoch_time_start = time.time()
9
10    def on_epoch_end(self, batch, logs={}):
11        self.times.append(time.time() - self.epoch_time_start)
12
13
14 early_stopping_callback = tf.keras.callbacks.EarlyStopping(
15     verbose=1,
16     patience=5,
17     restore_best_weights=True
18 )
19
20 reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, verbose=1,
21                               patience=3, min_lr=0)

```

```
1 METRICS = ['accuracy']
```

```

1 model.compile(
2     optimizer=tf.keras.optimizers.Adam(
3         learning_rate=1e-4
4     ),
5     # optimizer=tf.keras.optimizers.SGD(), # for VGG16_VGGFACE
6     loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
7     metrics=METRICS,
8 )

```

▼ Train Model

```

1 EPOCHS = 100
2 time_callback = TimeHistory()
3 history = model.fit(
4     x=train_ds,
5     validation_data=val_ds,
6     epochs=EPOCHS,
7     callbacks=[time_callback, early_stopping_callback, reduce_lr]
8 )

```

```

↺ Epoch 1/100
531/531 [=====] - 500s 898ms/step - loss: 1.2592 - accuracy: 0.6591 - val_loss: 2.1001 - val_accuracy: 0.47
Epoch 2/100
531/531 [=====] - 389s 732ms/step - loss: 0.2766 - accuracy: 0.9060 - val_loss: 3.1878 - val_accuracy: 0.47
Epoch 3/100
531/531 [=====] - 396s 745ms/step - loss: 0.1419 - accuracy: 0.9526 - val_loss: 3.3680 - val_accuracy: 0.47
Epoch 4/100
531/531 [=====] - ETA: 0s - loss: 0.0990 - accuracy: 0.9667
Epoch 4: ReduceLROnPlateau reducing learning rate to 1.9999999494757503e-05.
531/531 [=====] - 388s 730ms/step - loss: 0.0990 - accuracy: 0.9667 - val_loss: 3.8887 - val_accuracy: 0.47
Epoch 5/100
531/531 [=====] - 381s 718ms/step - loss: 0.0402 - accuracy: 0.9870 - val_loss: 4.0004 - val_accuracy: 0.47
Epoch 6/100
531/531 [=====] - ETA: 0s - loss: 0.0246 - accuracy: 0.9922Restoring model weights from the end of the best
531/531 [=====] - 381s 717ms/step - loss: 0.0246 - accuracy: 0.9922 - val_loss: 4.1271 - val_accuracy: 0.47
Epoch 6: early stopping

```

```
1 EPOCHS = len(time_callback.times)
```

```

1 # Create model path
2 model_path = Path(const.models_path, 'Frame', model._name)
3 run_folders = list(Path(const.models_path, 'Frame', model._name).glob('Run_*'))
4
5 if not run_folders:
6     model_path = Path(model_path, 'Run_1')
7 else:
8     last_run = run_folders.pop()
9     last_run_idx = Path(last_run).name.split('_')[-1]
10    model_path = Path(model_path, f'Run_{int(last_run_idx) + 1}')
11
12 model_path.mkdir(parents=True, exist_ok=False)

```

```

1 # Save info on the indexes used for train, val and test
2 ds_info_path = Path(model_path, f'{model._name}_dataset.txt')
3 with open(ds_info_path, 'w+', newline='') as res_file:
4     res_file.write(f'Train indexes: {train_idxs}\n')
5     res_file.write(f'Train files: {train_ds_elements}\n')
6     res_file.write(f'Val indexes: {val_idxs}\n')
7     res_file.write(f'Val files: {val_ds_elements}\n')
8     res_file.write(f'Test indexes: {test_idxs}\n')
9     res_file.write(f'Test files: {test_ds_elements}\n')

```

```

1 metrics = history.history

```

```


1 mod_loss = metrics['loss']
2 mod_val_loss = metrics['val_loss']
3 mod_accuracy = metrics['accuracy']
4 mod_val_accuracy = metrics['val_accuracy']
5 # mod_f1 = metrics['fBeta_score']
6 # mod_val_f1 = metrics['val_fBeta_score']
7
8 mod_mean_loss = np.mean(mod_loss)
9 mod_mean_val_loss = np.mean(mod_val_loss)
10 mod_mean_accuracy = np.mean(mod_accuracy)
11 mod_mean_val_accuracy = np.mean(mod_val_accuracy)
12 # mod_mean_f1 = np.mean(mod_f1)
13 # mod_mean_val_f1 = np.mean(mod_val_f1)

```

```

1 # Save Loss
2 plt.title('Model loss')
3 plt.plot(history.epoch, mod_loss, mod_val_loss)
4 plt.legend(['loss', 'val_loss'])
5 plt.ylim([0, max(plt.ylim())])
6 plt.xlabel('Epoch')
7 plt.ylabel('Loss [CrossEntropy]')
8 plt.savefig(Path(model_path, 'loss.png'))


```



```

1 # Save Accuracy
2 plt.title('Model accuracy')
3 plt.plot(
4     history.epoch,
5     100 * np.array(mod_accuracy),
6     100 * np.array(mod_val_accuracy)
7 )
8 plt.legend(['accuracy', 'val_accuracy'])
9 plt.ylim([0, 100])
10 plt.xlabel('Epoch')
11 plt.ylabel('Accuracy [%]')
12 plt.savefig(Path(model_path, 'accuracy.png'))

```



Unsupported Cell Type. Double-Click to inspect/edit the content.

Unsupported Cell Type. Double-Click to inspect/edit the content.

✓ Evaluate Model

```
1 model_eval = model.evaluate(test_ds, return_dict=True)
```

141/141 [=====] - 99s 701ms/step - loss: 1.7415 - accuracy: 0.4535

```
1 # Save model
2 model.save(Path(model_path, f'{model._name}.keras'), overwrite=False)
3 # Save history
4 np.save(Path(model_path, f'{model._name}_history.npy'), history)
5 # Save model image
6 model_img = tf.keras.utils.plot_model(
7     model, Path(model_path, f'{model._name}.png'), show_shapes=True,
8     show_layer_names=True, show_layer_activations=True
9 )
```

```
1 model_eval
```

{'loss': 1.7415306568145752, 'accuracy': 0.45348450541496277}

```
1 test_loss = model_eval['loss']
2 test_accuracy = model_eval['accuracy']
3 # test_f1 = model_eval['fBeta_score']
4 mean_epoch_time = np.mean(time_callback.times)
```

```
1 # Salvataggio informazioni modello
2 model_save_path = Path(model_path, f'{model._name}_result.txt')
3 with open(model_save_path, 'w+', newline='') as res_file:
4     res_file.write(f'BATCH: {BATCH_SIZE}\n')
5     res_file.write(f'Train loss: {str(mod_loss)}\n')
6     res_file.write(f'val_loss: {str(mod_val_loss)}\n')
7     res_file.write(f'Train accuracy: {str(mod_accuracy)}\n')
8     res_file.write(f'Train val_accuracy: {str(mod_val_accuracy)}\n')
9     # res_file.write(f'Train f1_score: {str(mod_f1)}\n')
10    # res_file.write(f'Train val_f1_score: {str(mod_val_f1)}\n')
11    res_file.write(f'Test loss: {str(test_loss)}\n')
12    res_file.write(f'Test accuracy: {str(test_accuracy)}\n')
13    # res_file.write(f'Test f1_score: {str(test_f1)}\n')
14    res_file.write(f'Mean epoch time: {str(mean_epoch_time)}')
```

```

1 # Salvataggio informazioni generali modelli
2 with open(Path(const.models_path, 'Frame', 'models.csv'), 'a+') as csvfile:
3     filewriter = csv.writer(
4         csvfile, delimiter=';', quotechar='|', quoting=csv.QUOTE_MINIMAL
5     )
6
7     # filewriter.writerow(
8     #     ["Model Name", "Epochs", "% Validation", "% Test set",
9     #     "Train loss", "Train accuracy", "Val loss", "Val accuracy",
10    #     "Test loss", "Test accuracy", "Mean epoch time", "Note"]
11    # )
12    test_ds_perc = utils.trunc((test_ds_elements * 100) / TOTAL_ELEMENTS, 2)
13    val_ds_perc = utils.trunc((val_ds_elements * 100) / TOTAL_ELEMENTS, 2)
14    full_path = str(Path(model_name, model_path.name))
15    filewriter.writerow(
16        [full_path, EPOCHS, val_ds_perc, test_ds_perc,
17         mod_loss, mod_accuracy, mod_val_loss, mod_val_accuracy,
18         test_loss, test_accuracy, mean_epoch_time, '']
19    )

```

```

1 for test_images, test_labels in test_ds.take(1):
2     print(test_images.shape)
3     print(test_labels.shape)

```

```

↔ (128, 224, 224, 3)
(128,)

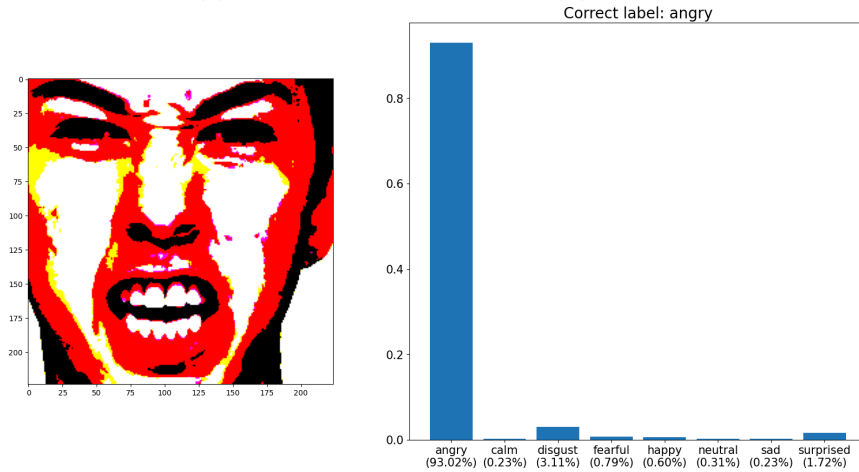
```

```

1 gen = np.random.default_rng(seed=None)
2 idx = gen.integers(0, len(test_images))
3 print(test_files[idx])
4
5 image = test_images[idx]
6 label = test_labels[idx]
7
8 net_input = utils.extend_tensor(image, 0)
9 prediction = model(net_input)
10 prediction = prediction[0].numpy()
11
12 valued_arr = []
13
14 for idx, name in enumerate(label_names):
15     valued_arr.append(f'{name}\n({prediction[idx]:.2%})')
16
17 fig, ax = plt.subplots(
18     nrows=1, ncols=2, width_ratios=[0.4, 0.6], figsize=(20, 10)
19 )
20
21 pltDisplay(image * 255, ax=ax[0]) # 1 channel
22 # pltDisplay(image, ax=ax[0])
23
24 ax[1].bar(valued_arr, prediction)
25 plt.xticks(fontsize=15)
26 plt.yticks(fontsize=15)
27 plt.title(f'Correct label: {label_names[label]}', fontsize=20)
28 # plt.xlabel('Predicted class')
29 # plt.ylabel('Percentage')
30 plt.show()

```

Generated/Frames/angry/21/01-01-05-02-01-01-21-035-01.jpg



Display a confusion matrix

Use a [confusion matrix](#) to check how well the model did classifying each of the commands in the test set:

Unsupported Cell Type. Double-Click to inspect/edit the content.

```
1 y_pred = model.predict(test_ds)
```

141/141 [=====] - 91s 643ms/step


```
1 y_pred = tf.argmax(y_pred, axis=1, output_type=tf.int32)
```

```
1 y_true = tf.concat(list(test_ds.map(lambda _, lab: lab)), axis=0)
```

```
1 report = classification_report(  
2     y_true, y_pred, target_names=label_names,  
3     output_dict=True, zero_division='warn'  
4 )
```

```
1 rep_to_csv = pd.DataFrame(data=report).transpose()
```

```
1 fig, ax = plt.subplots(figsize=(10, 10))
2 ax.xaxis.tick_top()
3 sns.heatmap(rep_to_csv.iloc[:NUM_CLASSES, :3],
4             cbar=True,
5             square=False,
6             annot=True,
7             annot_kws={'size': 15},
8             fmt='.2g',
9             linewidths=0.5)
10 plt.savefig(Path(model_path, f'{model._name}_f1_score.png'))
11 plt.show()
12
13
14 with sns.axes_style('white'):
15     fig, ax = plt.subplots(figsize=(3, 5))
16     ax.xaxis.tick_top()
17     sns.heatmap(rep_to_csv.iloc[:NUM_CLASSES, 3:],
18               cbar=False,
19               square=False,
20               annot=True,
21               annot_kws={'size': 15},
22               fmt='.4g',
23               cmap=ListedColormap(['purple'])),
24             linewidths=0.5)
25 plt.savefig(Path(model_path, f'{model._name}_support.png'))
26 plt.show()
```



```
1 report_save_path = Path(model_path, f'{model._name}_report.csv')
2 rep_to_csv.to_csv(report_save_path)
```

Unsupported Cell Type. Double-Click to inspect/edit the content.

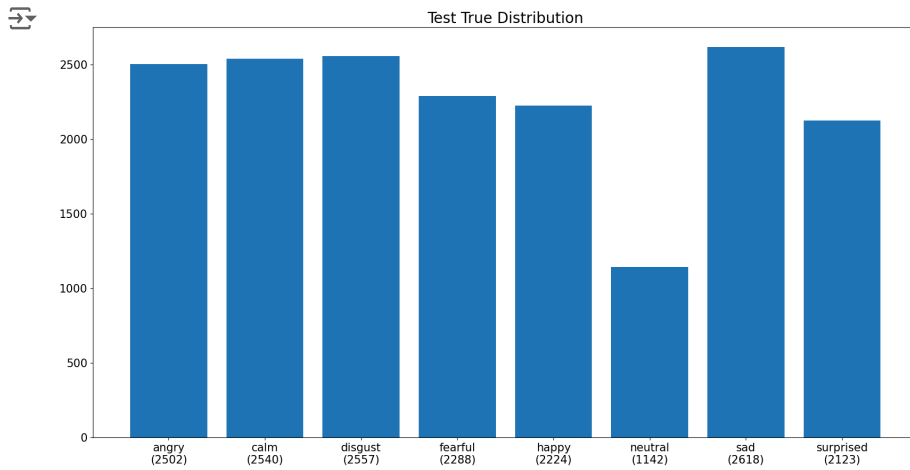
```
1 np.mean([report[c]['f1-score'] for c in list(report)[:NUM_CLASSES]])
```

0.442763100084397

```

1 # True Test Distribution
2 # unique, counts = np.unique(y_true, return_counts=True)
3 # collections.Counter(y_true)
4 counts = [np.count_nonzero(y_true == idx) for idx in range(len(label_names))]
5 valued_arr = []
6 # for i in range(len(label_names)):
7 for idx, name in enumerate(label_names):
8     count = counts[idx]
9     valued_arr.append(f'{name}\n({count})')
10
11 fig = plt.subplots(figsize=(20, 10))
12 plt.bar(valued_arr, counts)
13 plt.xticks(fontsize=15)
14 plt.yticks(fontsize=15)
15 plt.title('Test True Distribution', fontsize=20)
16 # plt.xlabel('Predicted class')
17 # plt.ylabel('Percentage')
18 plt.savefig(Path(model_path, f'{model._name}_trueDist.png'))
19 plt.show()


```




```

1 # Predicted Test Distribution
2 # unique, counts = np.unique(y_pred, return_counts=True)
3 counts = [np.count_nonzero(y_pred == idx) for idx in range(len(label_names))]
4 valued_arr = []
5 unique_idx = 0
6 for idx, name in enumerate(label_names):
7     count = counts[idx]
8     valued_arr.append(f'{name}\n({count})')
9
10 fig = plt.subplots(figsize=(20, 10))
11 plt.bar(valued_arr, counts)
12 plt.xticks(fontsize=15)
13 plt.yticks(fontsize=15)
14 plt.title('Test result distribution', fontsize=20)
15 # plt.xlabel('Predicted class')
16 # plt.ylabel('Percentage')
17 plt.savefig(Path(model_path, f'{model._name}_predDist.png'))
18 plt.show()

```




```

1 confusion_mtx = tf.math.confusion_matrix(y_true, y_pred)
2 fig, ax = plt.subplots(figsize=(10, 10))
3 sns.heatmap(confusion_mtx,
4             xticklabels=label_names,
5             yticklabels=label_names,
6             annot=True, fmt='.4g')
7 plt.xlabel('Prediction')
8 plt.ylabel('True Label')
9 plt.savefig(Path(model_path, f'{model._name}_heat.png'))
10 plt.show()

```

1}



```

1 def multiclass_roc_auc_score(target, y_test, y_pred, average="macro"):
2     # function for scoring roc auc score for multi-class
3     lb = LabelBinarizer()
4     lb.fit(y_test)
5     y_test = lb.transform(y_test)
6     y_pred = lb.transform(y_pred)
7
8     if len(target) > 2:
9         for (idx, c_label) in enumerate(target):
10
11             fpr, tpr, thresholds = roc_curve(
12                 y_test[:, idx].astype(int),
13                 y_pred[:, idx]
14             )
15             c_ax.plot(
16                 fpr, tpr, label='%s (AUC:%0.2f)' % (c_label, auc(fpr, tpr))
17             )
18     else:
19         fpr, tpr, thresholds = roc_curve(
20             y_test,
21             y_pred
22         )
23         c_ax.plot(
24             fpr, tpr, label='Model (AUC:%0.2f)' % (auc(fpr, tpr)), color='#ff7f0e'
25         )
26     c_ax.plot(fpr, fpr, color='b', linestyle='--', label='Random Guessing')
27
28     return roc_auc_score(y_test, y_pred, average=average)


```

```

1 # set plot figure size
2 fig, c_ax = plt.subplots(1, 1, figsize=(12, 8))
3
4 print('ROC AUC score:', multiclass_roc_auc_score(
5     label_names,
6     tf.reshape(y_true, (y_true.shape[0], 1)),
7     tf.reshape(y_pred, (y_pred.shape[0], 1))
8 ))
9
10 c_ax.legend()
11 c_ax.set_xlabel('False Positive Rate')
12 c_ax.set_ylabel('True Positive Rate')
13 plt.savefig(Path(model_path, f'{model1._name}_ROC.png'))
14 plt.show()

```

ROC AUC score: 0.6944431334629377



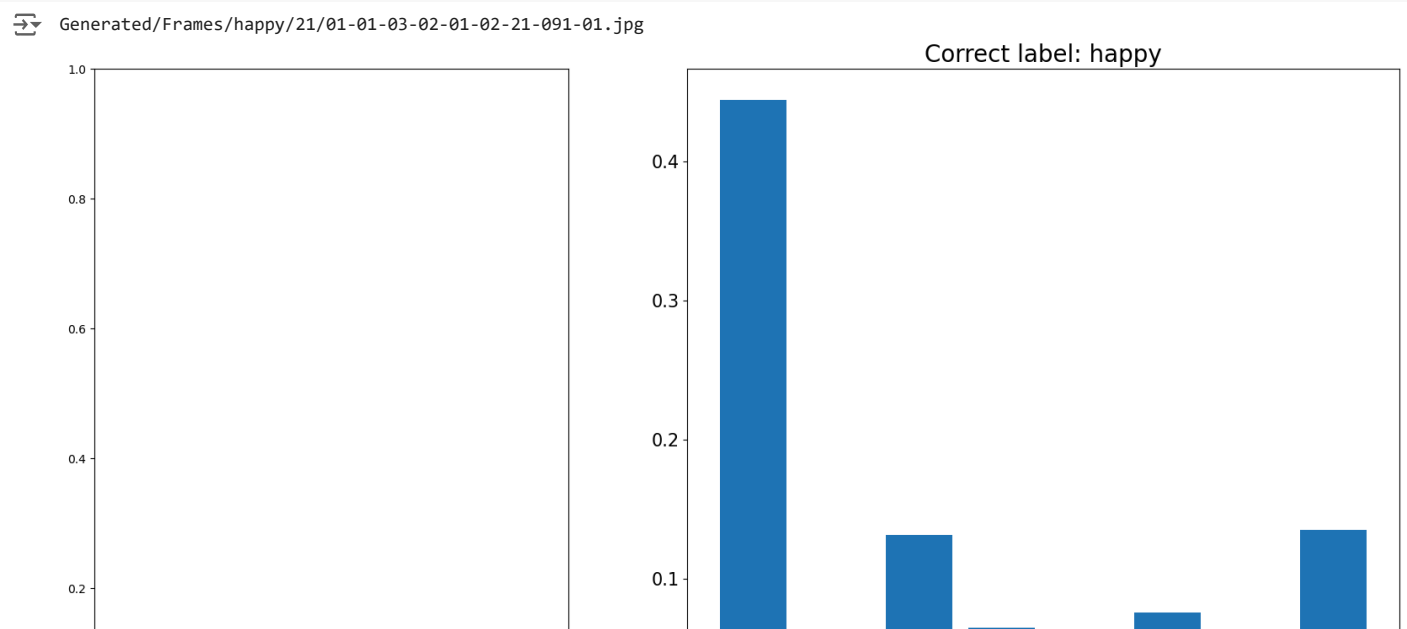
✓ Run inference on an image file

Finally, verify the model's prediction output using an image

```

1 # Take the first item of each class in test_files
2 # path = Path(next((subs for subs in test_files if 'angry' in subs), None))
3 # path = Path(next((subs for subs in test_files if 'calm' in subs), None))
4 # path = Path(next((subs for subs in test_files if 'disgust' in subs), None))
5 # path = Path(next((subs for subs in test_files if 'fearful' in subs), None))
6 path = Path(next((subs for subs in test_files if 'happy' in subs), None))
7 # path = Path(next((subs for subs in test_files if 'neutral' in subs), None))
8 # path = Path(next((subs for subs in test_files if 'sad' in subs), None))
9 # path = Path(next((subs for subs in test_files if 'surprised' in subs), None))
10 # path = Path('Generated/Frames/disgust/01-01-07-01-01-01-16-049-01.jpg')
11 print(path)
12 # print(tf.io.read_file(str(path)))
13 label = path.parent.parent.name
14
15
16 image = cv2.imread(str(path))
17 # image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
18 image = cv2.resize(image, (IMG_HEIGHT, IMG_WIDTH))
19 image = image / 255
20
21 net_input = utils.extend_tensor(image, 0)
22 prediction = model(net_input)
23 prediction = prediction[0].numpy()
24
25 valued_arr = []
26
27 for idx, name in enumerate(label_names):
28     valued_arr.append(f'{name}\n({prediction[idx]:.2%})')
29
30 fig, ax = plt.subplots(
31     nrows=1, ncols=2, width_ratios=[0.4, 0.6], figsize=(20, 10)
32 )
33
34 # pltDisplay(image * 255, ax=ax[0]) # 1 channel
35 # pltDisplay(image, ax=ax[0]) # 3 channels
36
37 ax[1].bar(valued_arr, prediction)
38 plt.xticks(fontsize=15)
39 plt.yticks(fontsize=15)
40 plt.title(f'Correct label: {label}', fontsize=20)
41 # plt.xlabel('Predicted class')
42 # plt.ylabel('Percentage')
43 plt.show()

```



✎ Export the model with preprocessing

0.0 0.2 0.4 0.6 0.8 1.0

The model's not very easy to use if you have to apply those preprocessing steps before passing data to the model for inference. So build an end-to-end version:

```

1 class ExportModel(tf.Module):
2     def __init__(self, model):
3         self.model = model
4

```